

- Linux Foundation: Hub for open-source tech. projects

- CNCF: Cloud Native Computing Foundation

under the Linux Foundation umbrella.

Vendor neutral home for many of the fastest growing projects on GitHub including Kubernetes

Prometheus → open source monitoring system and was the 2nd project accepted into CNCF after Kubernetes.

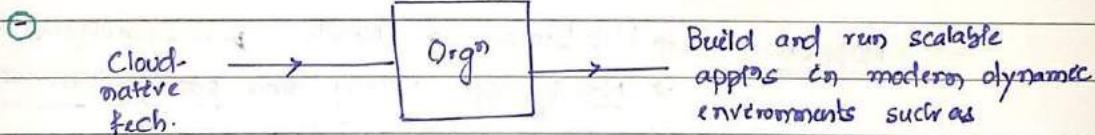
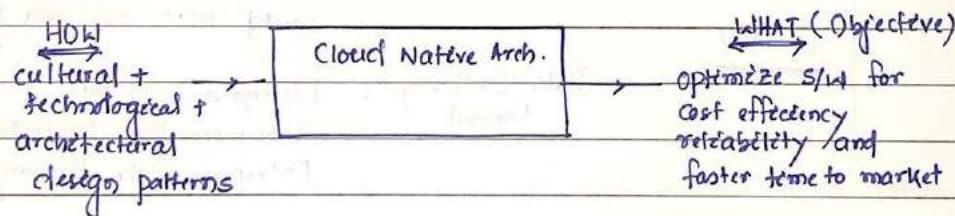
* code hosting platform for version control and collaboration

* commonly used to host open-source projects

Envoy → High-performance proxy developed in C++ to mediate all inbound and outbound traffic for all services in the service mesh.

- CLOUD NATIVE ARCHITECTURE:

① Cloud Native Arch. Fundamentals:



(Micro-service + CI/CD + Containers arch. + Cloud)

APIs +
↓
S/I intermediacy
that allows 2 apps talk to
each other.

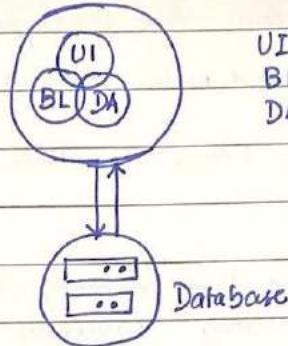
↓
Loosely coupled systems (resilient
manageable
observable)

e.g. Waiter → API
(tells Kitchen what the customer
wants and delivers after prepaⁿ of food)

↓
Robust customⁿ → frequent changes
predictably
with minⁿ fail

Monolithic Arch:

- * Traditional way of building apps



UI - User interface

BL - Biz. logic

DA - Data Access

- * a monolithic appⁿ is built as a single and undistributable unit.

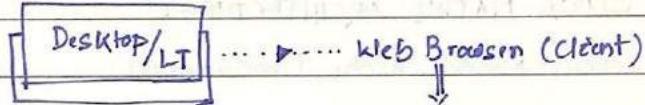
- * It is a single logical executable.

* S/W built on this model operates with one single base of code. Whenever stakeholders want to make updates or changes, they access the same set of code.

- * Monolithic meaning Object made from single, large piece of material

- * monolithic archr. also known as multi-tier archr.

- * Org's:



could HTML, CSS and Javascript

LEGACY

Data Exchange Format : Enterprise Data Bus
Enterprise Data Interchange
Enterprise Data Exchange

MODERN

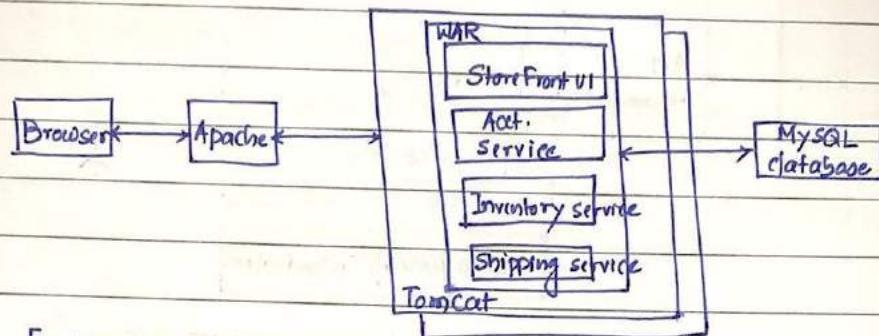
→ Web Browser → Mobile → SmartPhones

→ Mobile appⁿ taking more priority than web appⁿ

Org's → need to expose data using APIs.
because legacy data exchange formats are not compatible with mobile appⁿ.

→ Mobile appⁿ is nearly useless without internet connectivity and backend services.

- Traditional Web Appⁿ Arch:



E-commerce appⁿ:

takes order from customers
verifies inventory and available credit
and ships them.

as
developed Java-web appⁿ → consists of a single WAR file Runs on a web container such as Tomcat

- Advantages:

- * easier debugging and testing
- * simple to deploy
- * simple to develop

- Challenges:

* Understanding - When a monolithic appⁿ scales up, it becomes too complicated to understand.

* Hard to implement changes in a large and complex appⁿ with highly tight coupling

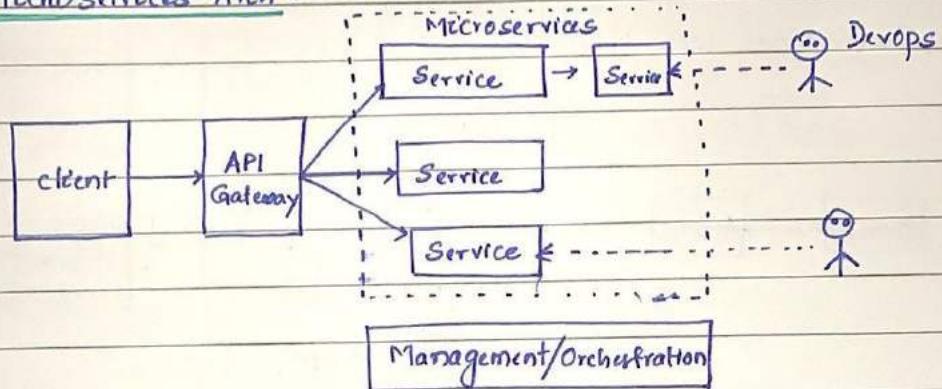
* code becomes complicated when the appⁿ grows.

* Scalability: cannot scale components independently, only the whole appⁿ.

* Extremely difficult to change techn. or language or framework as everything is tightly coupled and depends upon one another

* Reliability: Bug in any module can potentially bring down the entire appⁿ.

Microservices Arch.:



- * A micro-services archr. consists of a collⁿ of small, autonomous services.

Service → self-contained and should implement a single biz. capability within a bounded context.

→ independent and loosely coupled.

→ services communicate with each other by using well-defined APIs.

- * API Gateway: Entry point for clients

Clients call the API gateway, which forwards the call to the appropriate services on the back end.

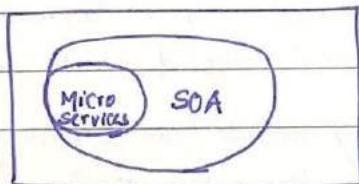
→ decouples clients from services.

★ API layer also enables the micro-services to communicate with each other over HTTP, gRPC and TCP/UDP.

- * The commⁿ between microservices is a stateless commⁿ

where each pair of request and response is independent.

- * Micro-services is an implementation of SOA. (service-oriented arch)



↓
S/W Components are exposed to the outer world for usage in the form of services.

SOA → Enterprise service exposure

Microservices → Applⁿ scope

- ★ Microservices-based arch. is aligned with event-driven arch. and SOA principles.
- ★ One of the greatest benefits of micro-services is scalability.

- REST: Representational State Transfer

framework for developing APIs

REST APIs work by sending HTTP requests and returning responses in JSON (JavaScript Object Notation) format.

* With HTTP methods (such as GET, POST, PUT and DELETE), clients can

- Principles of Microservice Design:

1. Reuse
2. Loose coupling
3. Autonomy
4. Fault tolerance
5. Composability
6. Discoverability

access and manipulate app's resources by using a URL.

Advantages:

- 1) Agility: easier to manage bug fixes and feature releases
 - update a service without redeploying the entire appⁿ
 - rollback an update if something goes wrong

2) Small code base

3) Mix of technologies

4) Fault isolation: If an individual micro-service becomes unavailable, it won't disrupt the entire appⁿ.

5) Scalability: Services can be scaled independently, letting you scale out subsystems that require more resources, without scaling out the entire appⁿ.

Drawbacks:

- 1) Additional complexity of creating a distributed system

→ developers must implement the inter-service commⁿ and deal with partial failure

→ implementing requests that spans multiple services is more difficult.

④ Transition from a legacy monolith to modern micro-services architecture:



- Running monoliths as micro-services (X)
 - Big-bang approach (Refactoring of the monolith, postponing the devp. of new features (X))
 - Incremental refactoring (New features - microservices, communicate with monolith through APIs
→ testing the interaction between services is difficult.)
- 2) Deployment complexity: operational complexity of deploying and managing a sys. comprising of many services
- 3) Increased memory consumption:

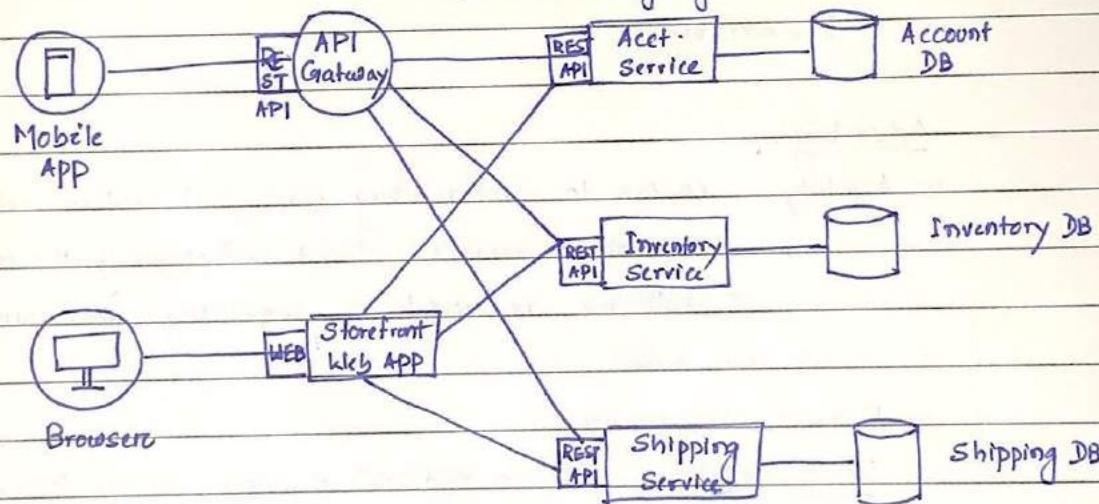
Monolithic appⁱⁿ → N instances

Microservice appⁱⁿ → NXM Service instances



overhead of M times as many JVM runtimes
(if each service runs in its own JVM)

- 4) Lack of governance: appⁱⁿ becomes hard to maintain well so many different languages and frameworks.



Monolithic Arch.

tightly
coupled services

(Changes in one module of code affect the other)

Internal procedures

(function calls)

comesⁱⁿ with
the appⁱⁿ

Decentralized data mgmt.

(entire app. may use one or
more DBs)

Micro-services Arch.

loosely
coupled services

(changes made in one doesn't
affect the other)

REST API calls

Decentralized
data mgmt.

(Each microservice
may use its own DB)



federated

Entire appⁿ must be scaled.

Each micro-service can be scaled independently.

Fault isolation - difficult
↳ If any specific feature is not working, the complete sys. goes down.

Fault isolation - easy



Characteristics of Cloud Native Architecture:

1. High level of automation → CI/CD pipeline → helps in easier DR regular automated
2. Self healing → health check of your appⁿ → automatic restart of services
3. Scalable → scaling your appⁿ describes the process of handling more load while producing the same cost. exp.
4. Cost-efficient
5. Easy to maintain
6. Secure by default

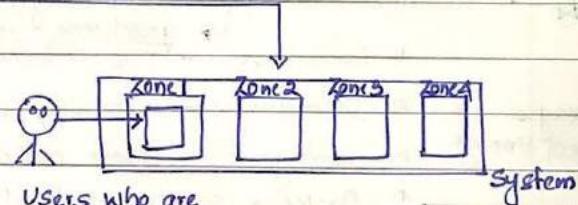
(Microservices)

One way of scaling can be starting multiple instances of the same appⁿ and distributing the load across them. Automating this process based on appⁿ metrics like CPU or memory can ensure availability and perf. of the services.

Scaling down your appⁿ and usage based pricing models (CSPs) can save cost if traffic is loco.

Kubernetes: optimization of infrastructure usage

efficient and denser placing of appⁿs



Zero-trust computing:

* known as zero-trust security model or zero-trust arch. (ZTA) or zero-trust network arch. (ZTNA)

* Approach to the design and implⁿ of IT systems where the main concept is 'never trust, always verify'.

↓
devices shouldn't be trusted by default, even if they are connected to LAN or managed network and previously verified.

* Zero trust is designed to protect modern digital envs by leveraging network segs, preventing lateral movement, providing layer 7 threat prevention and simplifying granular user-access control.

* Trust is a vulnerability.



* Security framework requiring all users, whether in or outside the org's network to be authenticated, authorized and continuously validated for security config and posture before being granted or keeping access to apps and data.

* Also known as perimeterless security.

Twelve-Factor App:

* Modern Era: S/W ^{delivered} as a service (Webapps or SaaS)

✳ * Twelve-Factor App: methodology for building software as a service.

↳ guideline for developing cloud-native apps

↳ Codebase - one codebase tracked in revision control (CI/CD)

Appn packages ← 1) Dependencies - explicitly declare and isolate dependencies
→ managed through 2) Config - store config's in an env. (no hard coding of config)
Sbt, maven 3) Backing Services - treat backing services as attached resources

(database, message broker or any ext. sys that the app communicates with) 4) Build, Release, Run - strictly separate build and run stages
5) Processes - Execute the app as one or more stateless processes
6) Port Binding - export services via port binding

7) Concurrency - scale out via the process model

8) Disposability - maximize robustness with fast startup and graceful shutdown

9) Dev/Prod parity - keep the envs as similar as possible

10) Logs - treat logs as event streams

11) Admin Processes - run admin tasks as one-off processes

↓
(Store the data in
store/DB instead of
in-memory)



- Concurrency: scaling the appⁿ

↓
running the appⁿ as multiple processes/instances instead of running as one large sys.

(scaling up) vertical scaling → add additional H/w to the sys. [e.g. CPU, RAM]

(scaling out) horizontal scaling → add additional instances of the appⁿ
(adoption of containers) → adding more machines to your pool of resources.



- System Desposability: sys. should not be impacted when new instances are added or taken off the existing instances as per need.

- Logs: observability: → achieved by APM tools (ELK) log aggr. tools (splunk)

- Auto Scaling:

* Auto Scaling: cloud computing tech. for dynamically allocating computational resources. (compute, memory or networking resources)

- Instance: single server or machine

- Advantages:

cost

Security

availability

Resilience

* Load Balancer: distribute the incoming traffic across multiple targets (VM instances)

↓
scales your app

supports heavy traffic

detect and automatically remove unhealthy VM instances

route traffic to the closest VM.



* Horizontal Auto Scaling: (Process of spawning new compute resources)

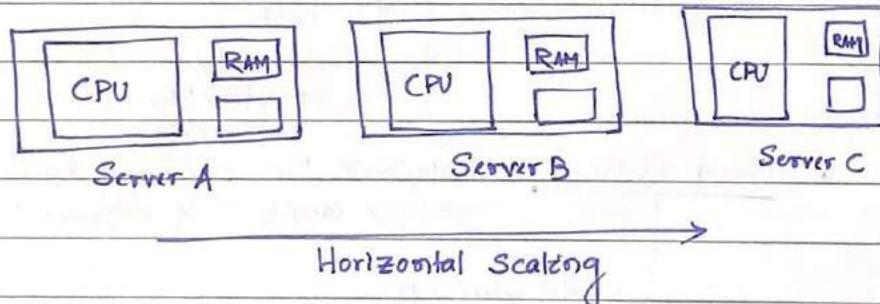
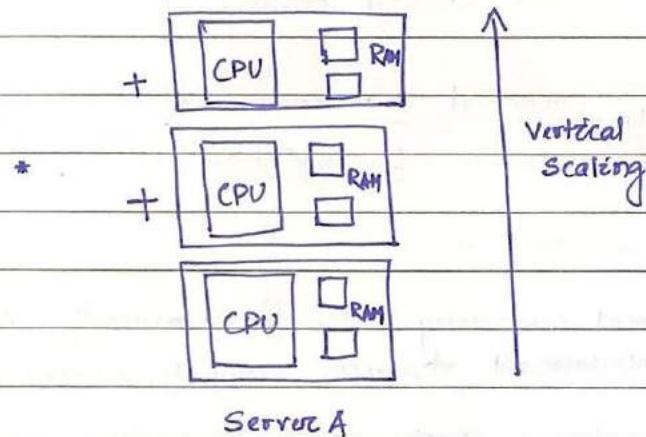
↓
new copies of your appⁿ process
VMs
new racks of servers
new servers or machines

* Vertical Auto Scaling: describes the change in size of underlying H/W which only works (CPU, memory) without hardware limits for BareMetal servers, VMs

The H/W can be scaled up e.g. by adding more RAM but only until all RAM slots are occupied.

* Phy. Server dedicated to a single tenant (single customer)

* Known as managed dedicated servers



- Serverless:

* Serverless computing \rightarrow arch. where code execution is managed by a cloud provider instead of the traditional method of developing apps and deploying them on servers.



* Developers don't have to worry about managing, provisioning and maintaining servers when deploying code.

- * Serverless relies on functions or more specifically functions-as-a-service (FaaS) where developers break down their app's into small, stateless chunks, meaning that they can execute without any context regarding the underlying server.

FaaS offerings:

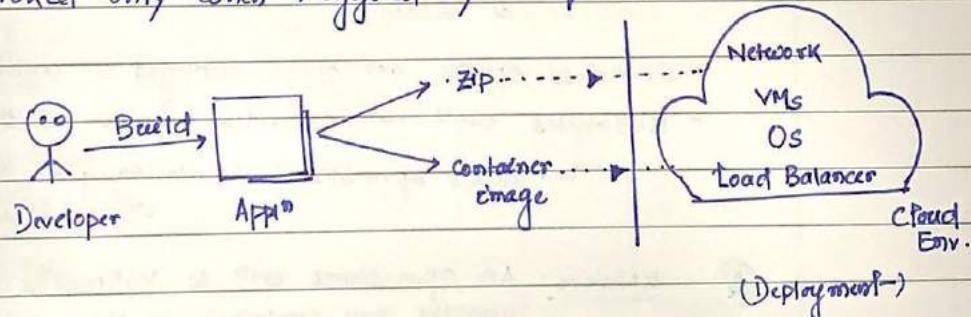
(actually a subset of Serverless)

- AWS Lambda : Allows developers to run code without provisioning or managing servers. AWS charges you for the compute power along with the storage you use.
- IBM OpenWhisk
- Cloud Run (Google)
- MS Azure Functions

- * FaaS $\xrightarrow[\text{also known as}]{}$ Caas (Compute as a Service)

- *  Developers $\xrightarrow[\text{don't have to know}]{}$ anything about the H/W or OS where the code runs on. It's taken care by the CSP.

- * Serverless functions are event-driven, meaning the code is invoked only when triggered by a request.



- * Advantages:

- 1) Lower costs
- 2) Simplified scalability
- 3) Simplified backend code
- 4) Quick turnaround

- * Drawbacks:
 - 1) Stable or predictable workloads
 - 2) Monitoring and debugging
 - 3) Vendor lock-in



- * Serverless offerings from public cloud providers are usually metered on demand through an event-driven execution model.

event-driven arch.
is a S/W arch.
and model for
app design.

Event: Any significant
occurrence or
change in state
for system H/W or S/W.

events can
generate from a user
like a
mouse click or keystroke
external source.

- * Kubernetes: Open-source container orchestration platform that automates schedules the deployment management and scaling of containers.

- Serverless apps are often deployed in containers.
- Kubernetes can't run serverless apps on its own.
 - needs specialized S/W to integrate Kubernetes well with CSP's serverless platform



- KNative: An open-source extn to Kubernetes that enables any container to run as a serverless workload on any cloud platform that runs Kubernetes.

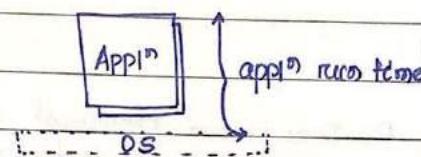


- Goal: Standardize how event publishers describes their events.
- CloudEvents: specify for describing event data in a common way.
 - to provide interoperability across services
 - one-type of event driven format
 - platforms and systems



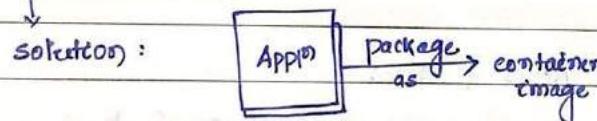
Events bases for scaling serverless workloads or triggering corresponding functⁿs.

- Open Standards:



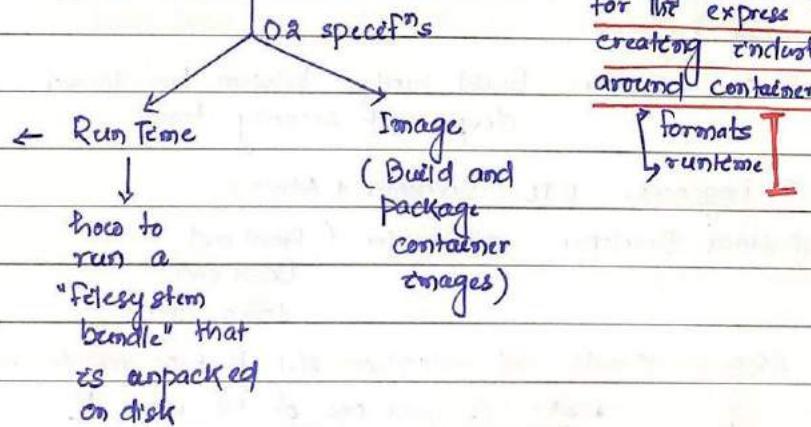
Common problem: Building and distribution of SW packages.

apps have a lot of req and dependencies for the underlying OS and the app1 run time.



OCI: (Open Container Initiative) → open governance structure for the express purpose of creating industry standards around container

(config), execution env. and life-cycle of containers



OCI impl → download an OCI image

unpack that image into an OCI Run-time filesystem bundle

OCI Run-time bundle to be run by OCI Runtime

OCI Distribution spec (API protocol) → defines distribution of content and container images.

CNI (Container Network Interface) → implement networking for containers
||
(CNCF proj.)

CRI (Container Runtime Interface) → implement container runtimes in container orchestrator systems

CSI (Container Storage Interface) → implement storage in container orchestrator systems

SMI (Service mesh interface) → implement service meshes in container orchestrator systems

Cloud Native Roles & SRE:

Cloud Architect: responsible for adoption of cloud

DevOps Engineer: CI/CD

Security Engineer:

DevSecOps Engineer: Build bridges between traditional devops and security teams

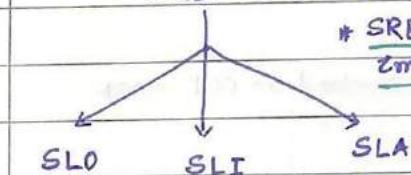
Data Engineer: ETL (Developers + Admin)

Full-Stack Developer: all-rounder (Front end
Back end
Infra. opns.)

SRE: Create and maintain software that is reliable and scalable.

* SRE is just one of the ways to implement the DevOps culture.

★ class SRE implements DevOps { }



↓
specify a target level for the reliability of the service

↳ a carefully defined quantitative measure of some aspect of the level of the service that is provided

(or time)

Error Budget: defines the amount of errors your appⁿ can have, before actⁿs are taken, like stopping deployments to prodⁿ.

Community and Governance:

- * Open Source (supported by CNCF)



SandBox Incubation Graduation

- * CNCF → Technical Oversight Committee

defining and maintaining tech. vision
approving new proj.

accepting feedback from end-user committee
defining common practices

practices the principle of
minⁿ viable governance

- * An SRE Team is responsible for the availability

latency

performance

efficiency

change mgmt.

monitoring

emergency response

capacity planning

Services

- * Core Tenets of SRE: durable focus on engg.

Pursuing maxⁿ change velocity without violating
monitoring (alerts, tickets, logging) or service's std

Emergency Response

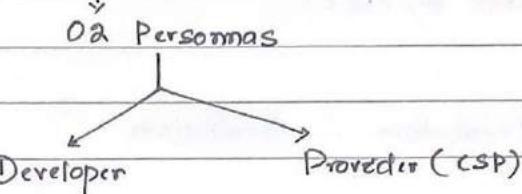
Change Mgmt.

Demand forecasting and capacity planning

Provisioning

Efficiency and Perf.

- Serverless: no upfront provisioning
no management of servers
Pay what you use for building apps



Backend as a service (BaaS):

Cloud service model on which the developers outsource all the behind the scenes aspects of a web/mobile app so that they have to maintain the frontend.

3rd party API based services that replace core subsets of functionality in an app

Best Practices for designing a Micro-Services Arch:

- 1) create a separate database for each microservice
- 2) keep code at similar level of maturity
- 3) do a separate build for each micro-service
- 4) deploy in containers
- 5) Treat servers as stateless

Principles for cloud-native arch.:

- 1) Design for automation
- 2) Be smart with state (design components to be stateless)
- 3) Favor managed services
- 4) Practice defense in depth
- 5) Always be architecting

CLOUD NATIVE ARCHITECTURE

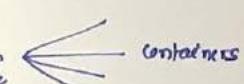
- Linux Foundation is a non-profit tech. consortium as a merger between Open Source Dev. Labs and the Free Standards Group to standardize Linux and promote its commercial adoption.
- CNCF:
 - * vendor neutral home for open source proj.
 - * 16 graduated proj.
 - * 26 incubating proj.
 - * 69 sandbox proj.
 - * Cloud native tech. empowers org's to build and run scalable apps in public, private and hybrid clouds. Containers, Service meshes, microservices, immutable infra. and declarative APIs exemplify this approach.
 - * Cloud native allows IT and S/W to move faster.
 - * Maintainers of CNCF hosted proj's manage the proj's.
- FaaS:
 - * is focussed on the event-driven computing
 - * valuable tool to efficiently and cost-effectively migrate apps to the cloud
 - * most central and most definitional element of the serverless stack
 - * kind of cloud computing service that allows developers to build, compute, run and manage app's packages as functions without having to maintain their own infra.
- KNative:
 - * It is an open-source proj. which extends K8s with serverless capabilities and a simplified dev. exp.
 - * KNative Serving defines a set of objects as K8s Custom Resource Definitions (CRDs). These obj's are used to define and control how your serverless workload behave on the cluster.
 - * open-source framework launched to bridge the gap between the containerized and serverless app's among cloud providers.
 - * serverless orch platforms.
- Open Container Initiative:
 - * The combi of the image manifest, image config and one or more filesystem serializations is called the OCI image.
 - * Image manifest → provides a config, and a set of layers for a single Container image for a specific arch. and OS.

- * OCI develops runc (container runtime) that implements their spec⁽¹⁾.

Container Network Interface :

- * CNI (Container Network Interface), a CNCF proj. consists of
- * CNI → concerned with the network connectivity of containers and releasing allocated resources when the container is deleted.
- * CNI → define a common interface between the network plugins and container execution.
- * K8s uses CNI as an interface between network providers and K8s pod networking.
- * CNI $\xrightarrow{\text{used by}}$ Container runtimes such as K8s, as well as Podman, CRI-O, Mesos and others.
- * The container/pod technically has no network interface. The container runtime calls the CNI plugin with verbs such as ADD, DEL, CHECK etc. ADD creates a new network interface for the container, and details of what is to be added are passed to CNI via JSON payload.

- Container Runtime Interface :

- ① * CRI is a plug-in which enables Kubelet to use a variety of container runtimes without the need to recompile.
- * CRI → Protocol Buffers + gRPC API + libraries
- * Kubelet $\xrightarrow{\text{communicates}}$ CRI Shim (S/I driver) $\xrightarrow{\text{implements the}}$ specific details of container engine
- * Docker doesn't support CRI.
- * Kubelet $\xrightarrow{\text{CRI protocol buffers}}$ CRI Shim \longrightarrow Container runtime 

- Container Storage Interface :

- * CSI enables K8s to work with a variety of storage devices.
- * Before CSI, volume plugins were ^{Serving} storage needs for container workloads in case of K8s.

②

- * CSI: benefit to the container orchestrators and storage vendors
- * Once you have deployed a CSI to a K8s cluster, it's available for use with

PV → persistent volume

Storage classes

PVC → persistent volume claim

- * CSI adopted by K8s

Clored Foundry

Mesos

Nomad

Service Mesh Interface:

- * A standard interface for service meshes on K8s.
- * defines a common standard that can be implemented by variety of service mesh providers.

④ * SMI enables flexibility and interoperability and covers the most common service mesh capabilities

- * Top 03 service mesh features:

- a) Traffic policy
- b) Traffic Telemetry
- c) Traffic mgmt.

- * Using the CNCF Envoy project, Open service mesh implements SMI for securing and managing your microservice apps.

SRE:

- * Treat opers as a S/W problem

- * 03 metrics:

SLA (commitment)

SLO (goals to meet)

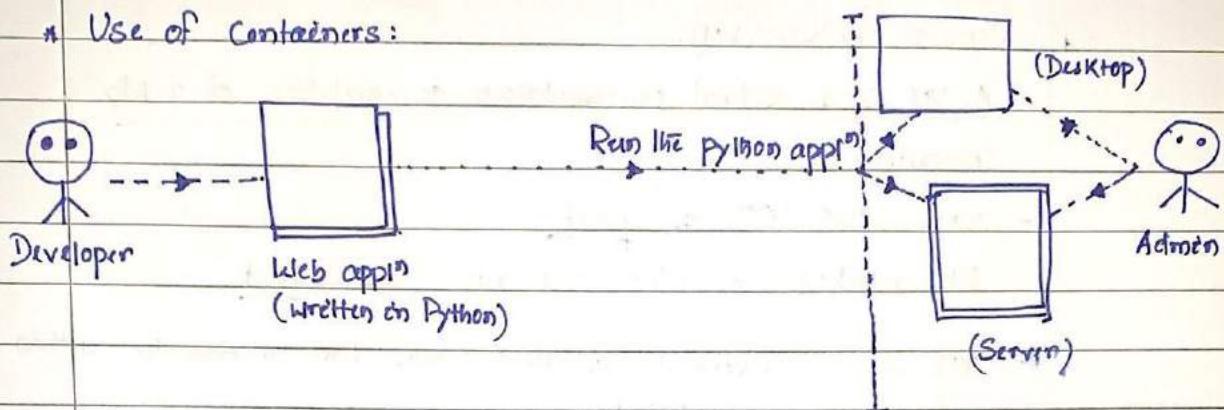
SLI (actuals)

- * SRE → automated solutions for operational aspects such as
on-call monitoring,
performance
capacity planning and
DR

- * Error Budget: The max^{able} amount of time a sys. can fail or underperform without violating the contractual terms of the SLA.

CONTAINER ORCHESTRATION:

- * Use of Containers:



Pre-requisites:

- 1) Install OS
- 2) Install the core Python packages to run the prog.
- 3) Install the Python ext's that the prog. uses
- 4) Configure networking for your system
- 5) connect to 3rd party systems like a database, cache or storage.

Developer: codes the app and knows its dependences

System Admin: Provides the infrastructure, installs all of its dependences and configures the system on which the app runs.

Drawback: process can be error-prone and hard to maintain

So, servers $\xrightarrow{\text{single purpose configured}}$ Running a database or an app server

- * To get more efficient use of server H/W, VMs can be used to emulate a full server with CPU

memory
storage
networking
OS and
S/W on top. \rightarrow This allocates running multiple isolated servers on the same H/W.

* Virtual Machine (VM):

- A VM is a virtual representation or emulation of a phy. computer.
- VM also known as guest.

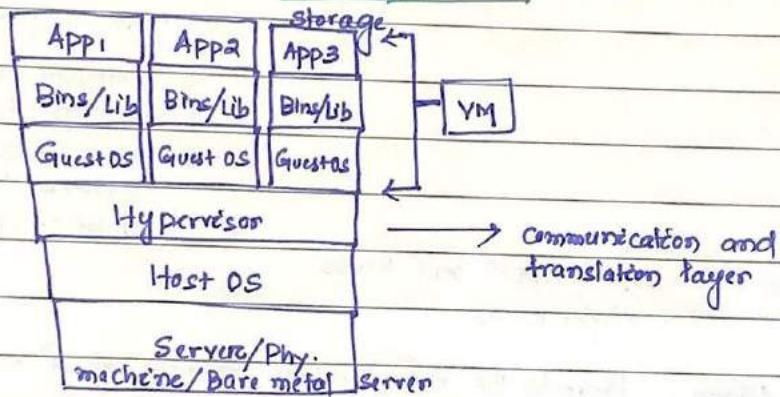
Phy. machines on which VMs run \rightarrow Host

- VM \rightarrow virtual env. that looks like a computer within a computer.

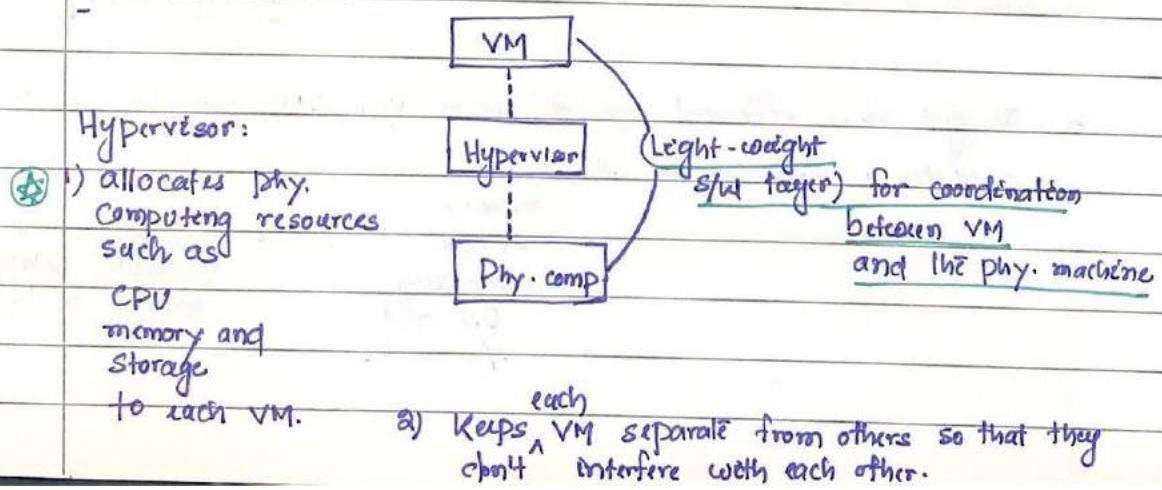
Has its own CPU

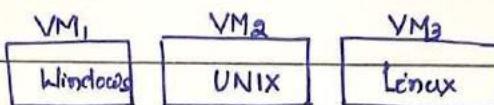
memory

network interface



- Each VM runs its own OS and functions separately from the other VMs even when they are running on the same host.
- VMs can be moved between host servers depending on demand.
- Virtualization makes it possible to create multiple VMs.



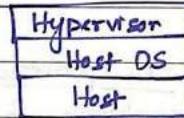


Hypervisor

Hypervisor:



- * Run directly on the phy. machine
- * Bare-metal hypervisors
- * Run on the host computer's OS
- * Hosted ~~hypervisors~~ hypervisors



- * e.g. Citrix Hypervisor
Microsoft HyperV
Oracle VM
- * Often found in server-based envs like enterprise data centers
- * e.g. VMware Workstation
VMware Player
- * Suitable for personal users or small biz.

The limit to how many VMs you can have are really limited by the amount of hard drive space.

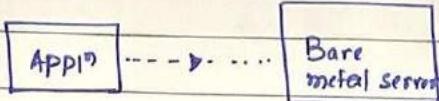
Advantages:

- 1) VMs can host multiple OS envs on a single phy. computer saving phy. space, time and mgmt. costs.
- 2) Agility and speed \rightarrow spinning up a VM is relatively easy and quick.
- 3) Lowered downtime \rightarrow portable and easy to move from one hypervisor to another

Disadvantages

- 1) VMs \rightarrow less efficient than real machines because they access the hw indirectly.
- 2) Unstable perf. if infra. req. are not met

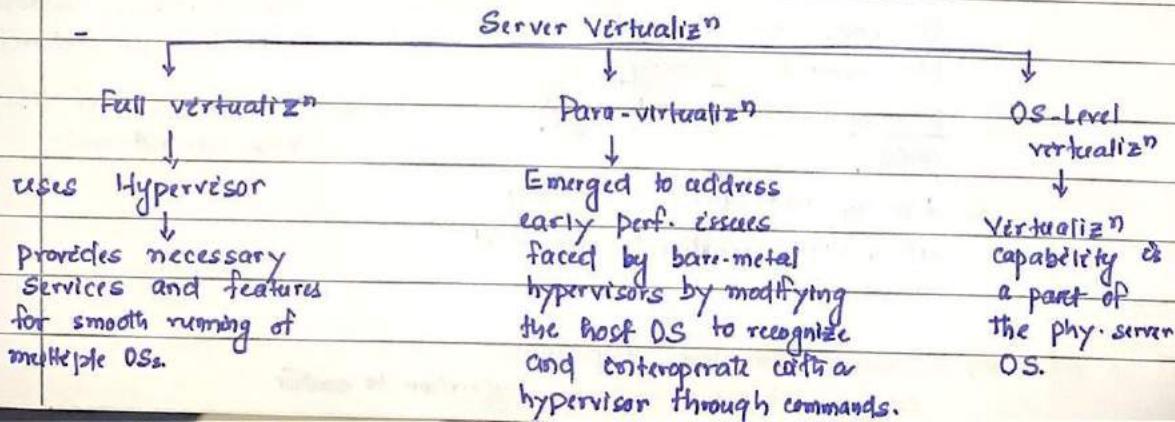
Why use a VM?



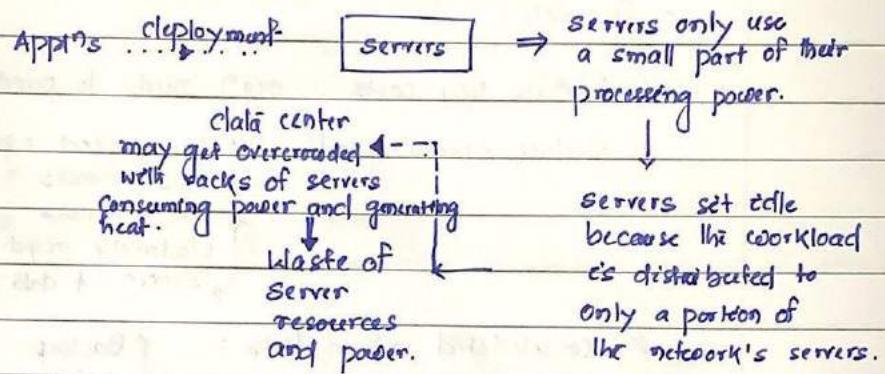
- * Most OS and app¹ deployments only use a small amount of the phy. resources available when deployed to bare metal.
- * By virtualizing your servers, you can place many virtual servers onto each phy. server to improve H/w utilization. This keeps you from needing to purchase additional phy. resources like hard drives or hard disks as well as reducing the need for power, space and cooling in the data center.
- * VMs provide additional disaster recovery options by enabling failover and redundancy that could previously only be achieved through additional H/w.
- * Because VMs are isolated, they are a good option for testing new apps or setting up a prod env.

Server Virtualization:

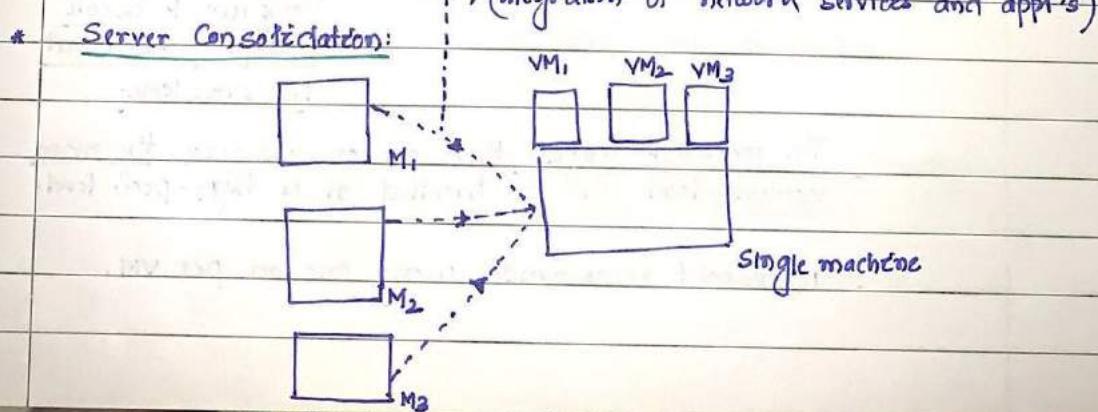
- Server Virtualizⁿ is the process of dividing a phy. server into multiple unique and isolated virtual servers by means of a S/W appl¹ called as hypervisor or VMM (virtual machine monitor).



- Benefits:
 - higher server ability
 - cheaper operating costs \rightarrow increased utilization of existing resources cuts down on the no. of servers needed.
 - eliminate server complexity
 - increased appn prof.
 - deploy workload quicker
 - improved disaster recovery \rightarrow capability to move data or appn's quickly from one server to another.
- Why server utilization?



- Drawbacks:
 - 1) perf. is less reliable. \Rightarrow Biz. critical appn's usually experience better perf. on bare metal servers compared to virtual ones.
 - 2) Capital investment is more expensive. \Rightarrow virtual cost of a virtualized server is usually more expensive than its bare metal counter part.
 - 3) Not all appn's can be virtualized. \Rightarrow some virtual envs must maintain some traditional server capacity and adapt into a hybrid ecosystem.



- Server consolidation: popular way to make more efficient use of H/W and resources.

↓

practice of reducing the number of servers or server func's in order to use computer resources more effectively and reduce costs.



- Virtualization $\xrightarrow{\text{enables}}$ server consolidation

- Benefits:

1) Reduce H/W costs : orgⁿ needs to purchase less servers.

2) Reduce operating costs :
 ↗ H/W support agreements
 O/S licenses
 S/W licenses
 ↗ electricity reqd. by server + data center

3) Consolidated admin. tasks :

↗ Backup
 patching
 monitoring
 Reporting

4) reduced environmental impact : Reducing the no. of phy. servers maintained is a huge step in reducing CO₂ emissions.

5) centralized mgmt.

* Virtual server consolidation Ratio : \rightarrow Number of virtual servers that can run on each phy. host machine.

VMware ESX 3.x \rightarrow 4 VMs/core
 Quad-core processor: 04 core \rightarrow presence of multiple cores allows your computer to execute multiple processes at the same time.

- The more CPU cores there are in a server, the more virtual-load can be handled at a high-perf. level.

- Microsoft recommends using one core per VM.

- Laptop: 2 cores
- Intel's 8th Gen: 4 cores
- CPU cores → made up of billions of microscopic transistors within a/ processor that help to make it work.
- How many VMs fit on a server? → It depends on your workload and it depends on how much phy. RAM you have.
- ④ People are worried that if a host has ~~too many~~ VMs, on it, perf. may degrade and the ^{failure of a} impact of the phy. host ^{on service delivery is worrisome.}
- ④ Hyper-threading: Process by which a CPU divides up its phy. cores into virtual cores that are treated as if they are actually phy. cores by the OS.
 - virtual cores called, threads
 - Intel's CPU (2 cores) → 4 threads or 4 virtual cores
- more threads means more work can be done in parallel.
- Increases CPU perf. by improving the processor's efficiency, thereby allowing you to run multiple demanding apps at the same time.
- c.g. Intel Pentium 4 (HT-enabled)
- Intel Pentium EE 840 (HT-enabled) processors support hyper-threading.
- RAM: Server's short-term memory
 - ↓
With more RAM, server can handle more tasks at once.
 - (16 GB - 6 TB)

Container Basics:

- * One of the earliest ancestors of modern container tech. is the chroot command.
 - ↓
isolate a process from the root filesystem
basically hide the files from the process
simulate a new root directory
 - changes the root directory for currently running processes as well as its child processes.

- * Root → refers → base directory (/)

chroot command → creates a fake env. inside which the process runs.

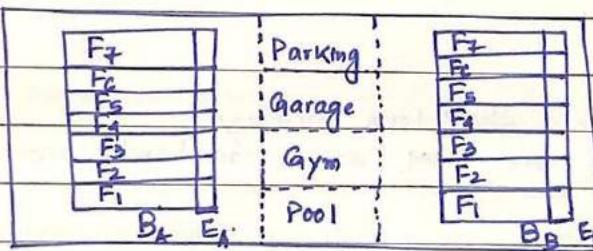
The processes run in such an env. cannot access files and commands outside that env.

or we can say that "the process is jailed within that env."

The created env. → called "chroot jail"

recovering your file system
reinstalling bootloader

- Namespace: → (Allows for isolation of global sys. resources between independent processes)
 - * Are a feature of the Linux Kernel that partitions kernel resources such that one set of processes sees one set of resources and another set of processes sees a different set of resources.



Building → B_A B_B
Elevators → E_A E_B

Part of same apartment society
Parking, Gym, Garage, Pool → Shared

↓
(computer)

(Apartment society)

(02 namespaces)

can reside on the same phy. computer can either share access to certain resources or can have exclusive access.

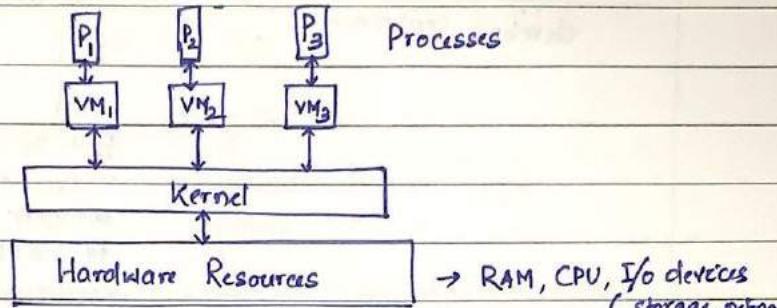
- * Linux Kernel 5.6 provides 8 namespaces:

1) pid 2) net 3) mount 4) ipc 5) user 6) uts 7) cgroup 8) time

containing
process
commands

Unix
time sharing

Linux Kernel:



- * Core interface between the computer's H/w and its processes.
- * Kernel is a computer program at the core of an operating sys. It is that part of OS that loads first and remains in the main memory.
- * Communicates between computer's H/w and its processes, managing resources as efficiently as possible.

Cgroups: (Control Groups)

- * is a Linux Kernel feature
- * Allows you to allocate resources (CPU time, system memory, network bandwidth or a comb) of these resources

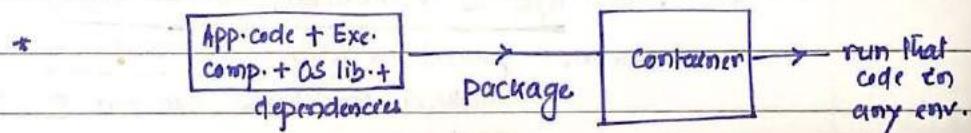
among user-defined group of tasks (processes)
running on a system.

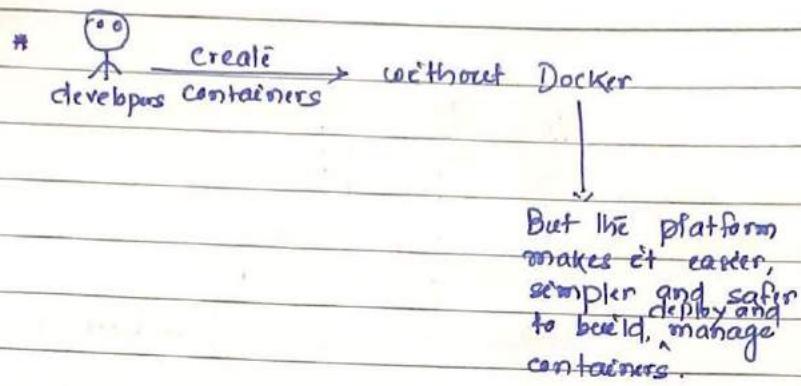
- * By using cgroups, sys. admins gain control over allocating, prioritizing, denying, managing and monitoring sys. resources

- * cgroups organized, hierarchically

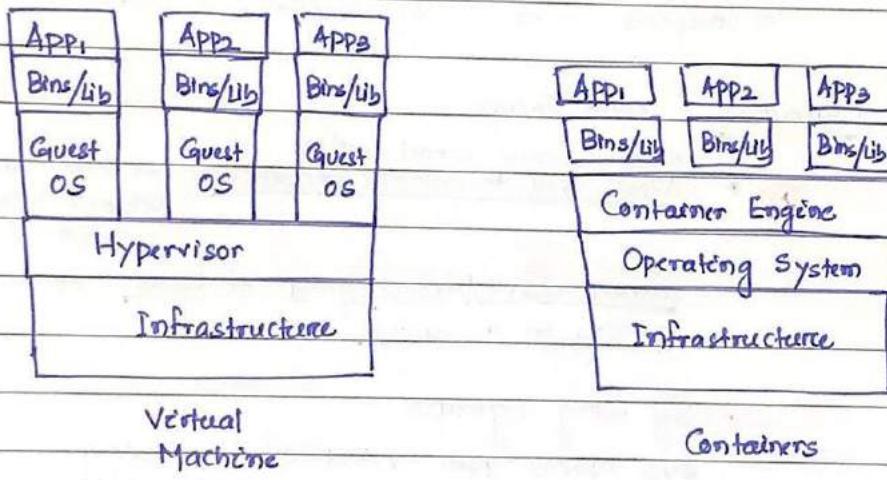
Docker:

- * Open-source platform for building, deploying and managing containerized apps





- * Docker released → public in 2013.
 - * Set of platform as a service (PaaS) products that use OS-level virtualization to deliver software in packages called containers.



- * VMs measured → GB
↓
contains their own OS, allowing them to perform multiple resource intensive funct's at once.

Containers $\xrightarrow{\text{measured}}$ MB

↓ package app source code, app¹⁰ dependencies, exe components and OS level libraries

- * VMs: It/w is virtualized to run multiple OS instances.

Containers: provide a way to virtualize an OS so that multiple workloads can run on a single OS instance.

* VMs: Machine-level isolation

Containers: Process isolation

* VMs: Hypervisor Creating the VMs.

Containers: Kernel namespaces
cgroups

resource
alloc

* VMs: Flexible in creating VMs by adjusting resource limits

Containers: Run your container image in any env.

It can be moved from one machine to another.

* VMs: VMware, KVM

Containers: Rancher

- Running Containers:

* To run industry-standard containers, you can follow the OCI run-time spec standard.

OCI (Open Container Initiative)

↓ maintains

a container runtime reference impl called runc.

* Container Run-time:

- also known as container engine
- S/w component that run containers on a host OS.
- e.g. runc, containerd, Docker, CRI-O

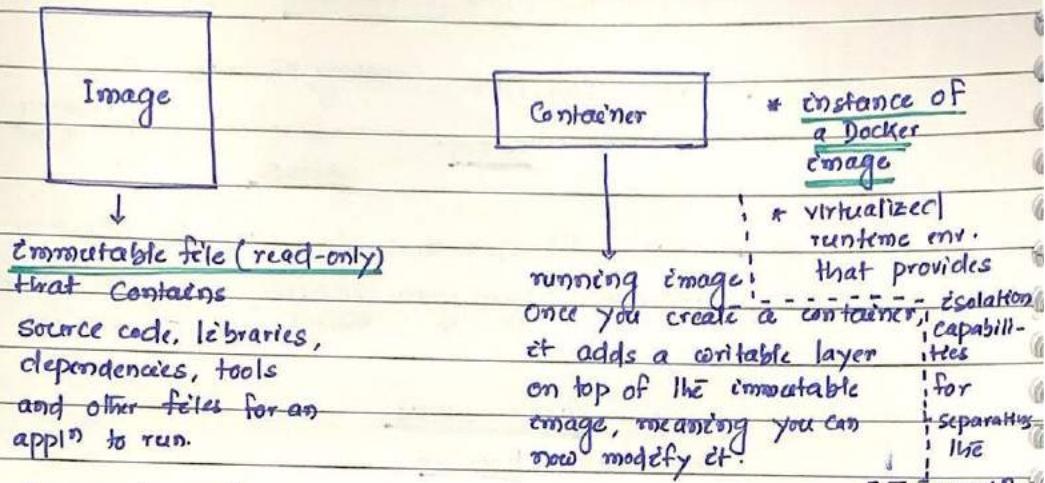
- CLI tool for spawning and running containers on Linux according to the OCI specifⁿ without having to run a Docker daemon.

- light-weight, portable
Container RunTime

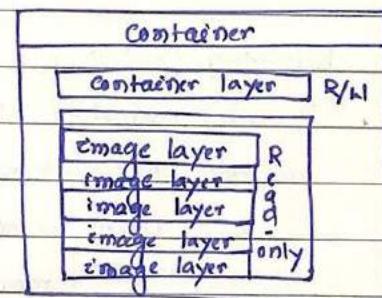
- runc is an open source project, written in Go lang.

- Host-level runtime

Relationship between container image and a container:



- * Docker image is made up of multiple layers, stacked to one another.



executed of app from the under-planting sys.

- * Images can exist without containers, whereas a container needs ^{to run} an image to exist. So, containers are dependent on images and use them to construct a run-time env. and run an app.



- * Image → logical entity
Containers → real world entity

- * Image → created only once
Containers → created any number of times using image.

- * Image → write script in Docker file
Container → Run "docker build" command

(creation)



- * Image → class
Container → instantiation of that class

→ OOPs concept

- Create images from a Dockerfile:

↓
Plain text file that contains
instr's that tell the Docker
Build engine how to create an image.

Run image:

docker run -it ubuntu:18.04 /bin/bash
image name

- Buildah: tool for building OCI-compatible images through a lower-level coreutils interface.

- Kaniko: tool to build container images from a Docker file, inside a container or Kubernetes cluster.

- Podman: (Pod Manager) → tool for managing OCI containers and pods

Container engine that's compatible
with the OCI containers spec'n.

Running containers without root privileges

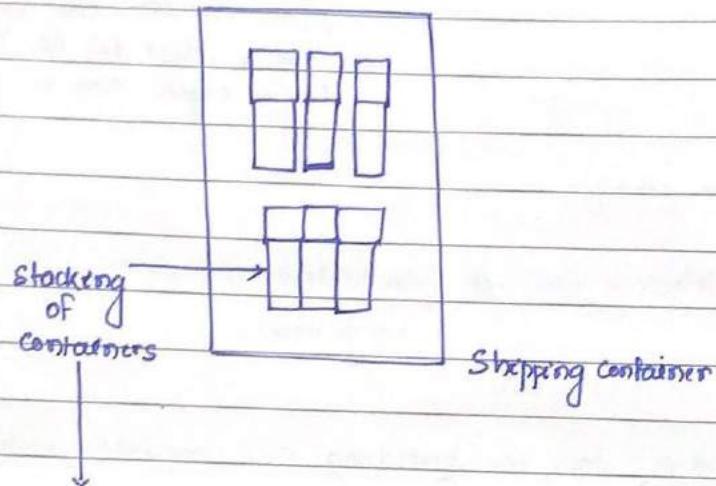
nginx → web server
docker run nginx:1.20 → run the container in foreground

docker run --detach --publish-all nginx:1.20 → run the container
in background mode
container id

docker ps → overview of all the running containers
docker stop <container id> → stop the container

Podman run --detach --publish-all nginx:1.20 → run the container
using podman

- Building Container Images:



easy to unload with a crane onto a truck
irrespective of what is inside

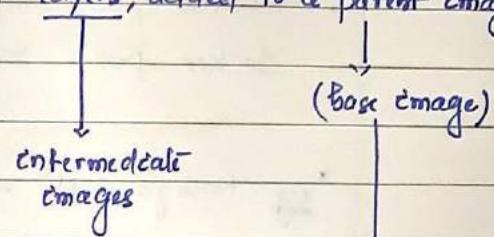
- Docker container: light-weight image
standalone executable
executable
- { package of s/w
that includes
everything to run an
appn:
code
runtime
sys. tools
sys. libraries and
settings}

- 2015: Docker image format → donated → OCI-image-spec

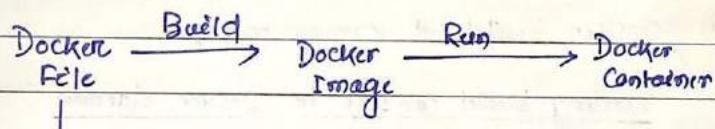
- Container images include everything a container needs to run.

Container engine - Docker or CoreOS
Sys. libraries
utilizers
config. settings and
specific workloads

Container image is composed of layers, added to a parent image.



1st layer of the Docker image
Found upon which all other layers are built
available in the Docker Hub (public
Container
repository)



Each entry in the Docker file adds a layer to the container image.

Building : docker build -t my-python-image -f Dockerfile

Container Registry: repository or coll^o of repositories used to store Container images for Kubernetes, DevOps and container based applic^o devp.

Server-side applic^o that stores and lets you distribute Docker images.

JFrog Container Registry → comprehensive registry
supporting Docker containers and
Helm Chart repositories

(upload and download
images)

* docker push my-registry.com/my-python-image
docker pull my-registry.com/my-python-image

* e.g. Node.js app → source code

i) Create a Dockerfile: → contains

```
FROM <base-image>
RUN <dependencies>
WORKDIR <setting the dir.>
COPY <copy the source code
      to the container>
RUN <install dependencies>
CMD <docker run>
```

ii) docker build & <image name>

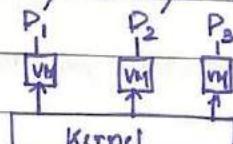
→ sending build context to Docker daemon

iii) docker images → lists all the images currently installed on the system

iv) docker run --detach --publish 3000:3000 <image name>
→ run the container

Security:

* When containers are started on the same machine, they always share the same kernel.

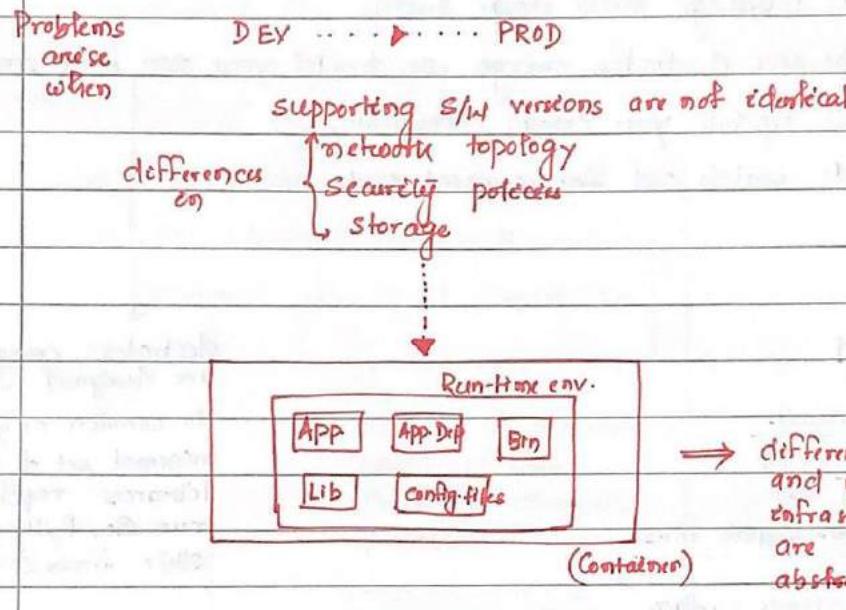
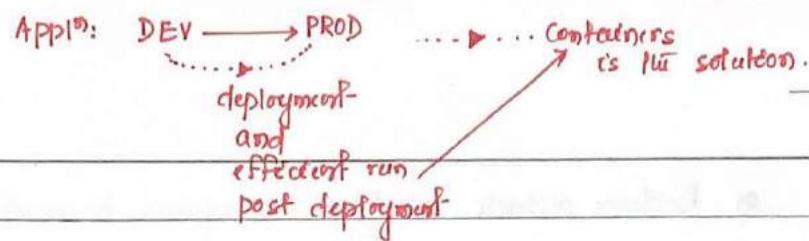


* If containers are allowed to call kernel functions to

kill other processes or modifying the host network by creating routing rules

(Risk to the whole system)

WHY CONTAINERS?



- * Starting processes as root or administrator → Security risk
- * Use of public images

Public image registries: Docker Hub
Quay

Risk

These images were not modified to include malicious SW.

Dockerfile Best Practices: (image builds)

1) Avoid unnecessary privileges:

a) avoid running containers as root (UID 0)

← b) don't bind to a specific UID (0644)

c) make executables owned by root and not writable

run the
Container as
non-root user

↓
root
owned by user
executed by non-root user
shouldn't be world writable

2) Reduce attack surface

- a) Leverage multi-stage builds
- b) use distroless images or build your own from scratch
- c) update your images frequently
- d) watch out for exposed ports

Multi-stage build
includes only the
minimal reqd. binaries
and dependencies in the
final image, and not
build tools or intermediate files.

reduces the attack surface
decreasing vulnerabilities

distroless images
are designed
to contain only the
minimal set of
libraries reqd. to
run Go, Python or
other frameworks.

3) Prevent confidential data leaks

- (use Docker secrets
Kubernetes secrets)
- a) never put secrets or credentials in Dockerfile instructions
 - b) prefer COPY over ADD
 - c) be aware of the Docker context, and use `.dockerignore`

"." parameters → build context
dangerous as you can copy confidential
or unnecessary files into the container.

4) Others

- (grouping multiple commands together will reduce the no. of layers)
- a) reduce the no. of layers and order them intelligently
 - b) add metadata and labels (image management-like including the app version)
 - c) leverage printers to customize checks
 - d) scan your images locally during devp.

Shift left security

Tools like Haskell Dockerfile linter (Hadolint) can detect
bad practices in your Docker file
and even expose issues inside the
Shell Commands executed by the RUN cmd.

5. Beyond image building

a) protect the clocker socket and TCP connections

b) sign your images and verify on runtime

c) avoid img mutability

d) don't run your env. as root

e) include a health check

f) restrict your app's capabilities

④ docker content trust, docker notary, harbor notary

digitally sign images

include a HEALTHCHECK cmd in your Dockerfile whenever possible

(Kubernetes → use liveness Probe config)

big privileged door into your host sys.

/var/run/clockr.sock

Correct permissions

In execution, restrict the app's capabilities to the minimal reqd. set

Tags are a volatile ref. to a concrete image version in a specific point in time.

--cap-drop flag (Docker)

Security Context. capabilities.drop (Kubernetes)

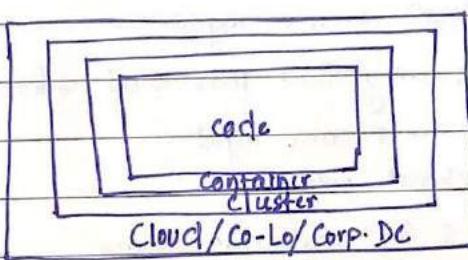
Cloud Native Security (4 C's)

Cloud

Cluster

Container

Code



Cloud: trusted computing base of a Kubernetes cluster.
CSPs → security best practices

cluster: infrastructure security

i) all access to Kubernetes control plane

→ not allowed publicly over the internet

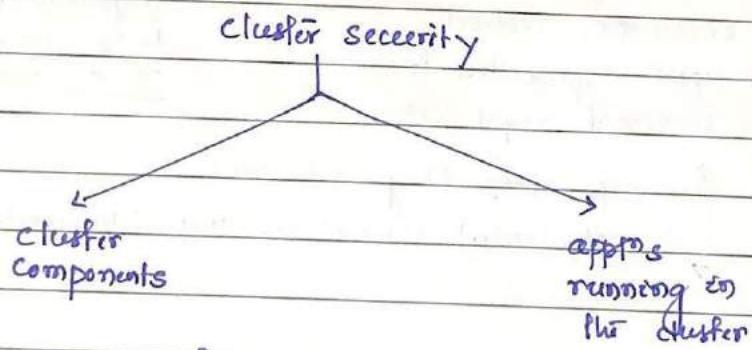
→ controlled by networks ACLs restricted
to the set of IP addrs needed to
administer the cluster.

ii) nodes configured to only accept conn's from
the control plane on the specified ports

iii) Kubernetes access to etcd API → principle of
least privilege

iv) access to etcd → should be limited to
control plane only

v) etcd encryption → etcd disks should be
encrypted at rest.



area of concern for workload security:

- 1) RBAC authorization → access to Kubernetes API
- 2) Authentication
- 3) API secrets mgmt.
- 4) ensuring that pods meet defined Pod Security Standards
- 5) cluster resource mgmt.
- 6) network policies
- 7) TLS for Kubernetes ingress

Container:

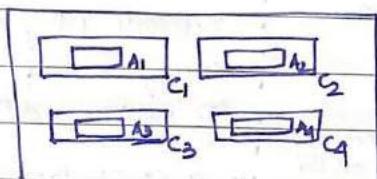
Container vulnerability scanning
image signing and enforcement
Disallow privileged users

use Container runtime with stronger isolation

Code:

access over TLS only
limit port ranges of comm
3rd party dependency security
static code analysis
dynamic probing attack → few automated tools
that you can run
against your service
to try some of the
well-known service
attacks.

- Container Orchestration Fundamentals:



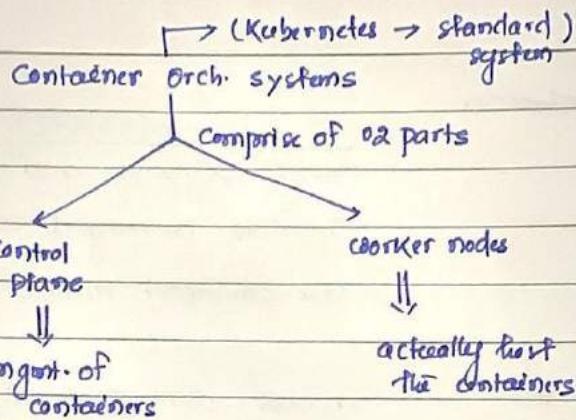
loosely coupled
isolated
independent

Services ⇒ micro-service arch.

App¹⁰

Container orch¹⁰ systems: → (cluster of multiple servers and host
Containers on top)

- 1) provide compute resources like VMs where containers can run on
- 2) schedule containers to servers in an efficient way
- 3) allocate resources like CPU and memory to containers
- 4) manage the containers and replace them if they fail
- 5) Scale containers if load increases
- 6) provision networking to connect containers together
- 7) provision storage of containers need to persist data



- Networking:

- * Microservice arch. depends on network comm.
- * Network namespace allows each container to have its own unique IP address.
- * Modern impls of container networking are based on the Container Network Interface (CNI).



- * The type of network a container uses is transparent from within the container.

↓
you need

Bridge → (when multiple containers need to communicate on the same Docker host)

Overlay Macvlan

Custom network plugin

Host

when you need.
Containers running on diff. Docker hosts to communicate

- * From a container's point of view, it has

Network interface with IP add.

Gateway

Routing table

DNS services

other net. details

- * Host network - when the network stack should not be isolated from the Docker host

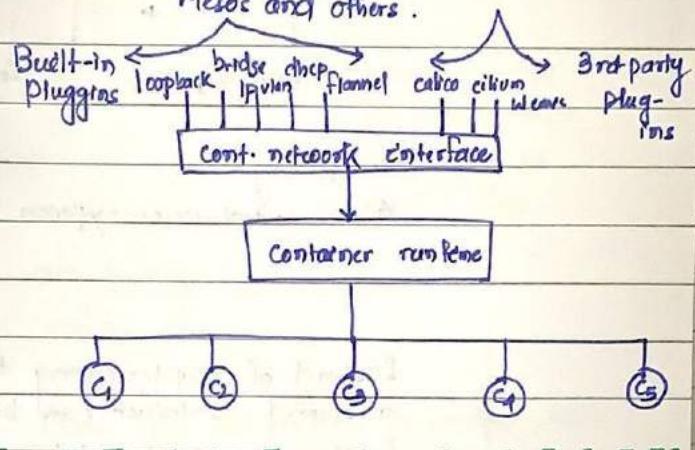
- CNI: standard that can be used to write or configure network plugins and makes it easy to swap out diff. plugins to various orch. platforms.

- a CNCF project consists of a specif. and libraries for writing plugins to configure network interfaces in Linux containers, along with a no. of plugins.

- CNI: deals with network connectivity of containers and removing allocated resources when the container is deleted.

- Kubernetes use CNI as an interface between network providers and Kubernetes pod networking.

- CNI is used by container runtimes such as Kubernetes, as well as Podman, CRI-O, Mesos and others.



Service Discovery & DNS:

* In container orchestrators,

- 100's/1000's of containers with individual IP addresses
- containers deployed on variety of diff. hosts, diff. data centers or even geo-locs
- containers or services need DNS to communicate. Using IP add. is nearly impossible.

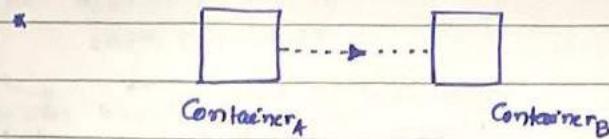
solution → automation → all info put in Service Registry

* Finding other services in the network and requesting info about them → Service discovery

*  Provide a Service API register new services as they are created

* Key-value store: store info about services
e.g. etcd, Consul, Apache Zookeeper

- Service Mesh:

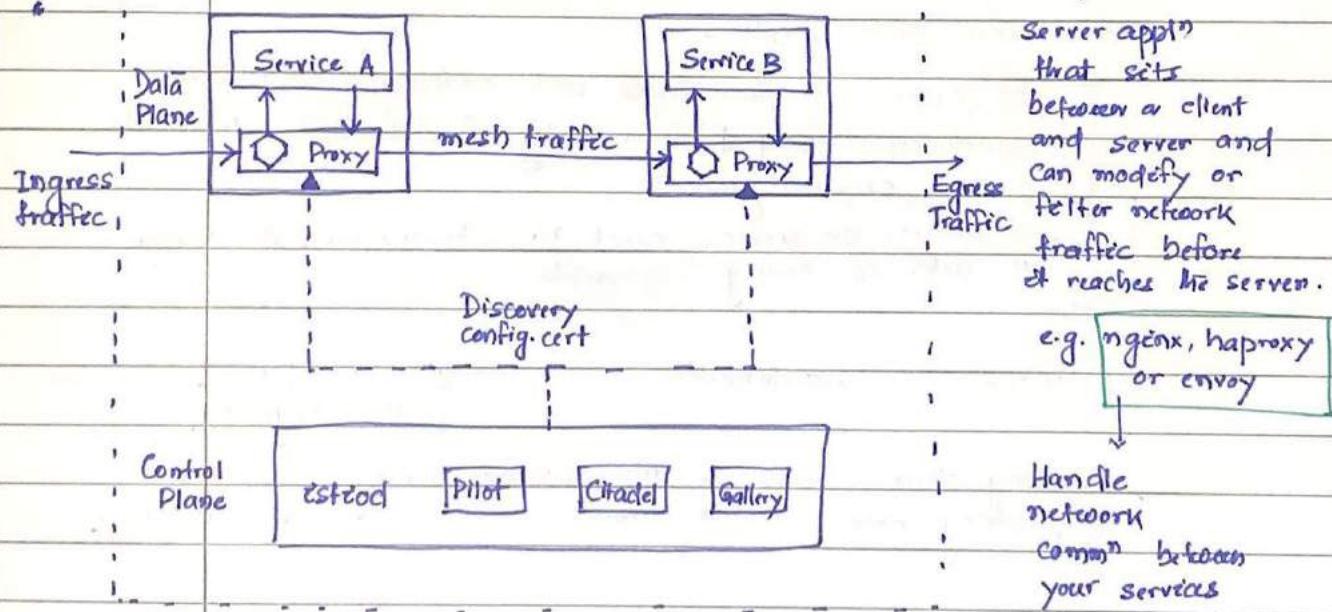


Monitoring
Access control or encryption

→ of network traffic clustered when
when
Containers
communicate
with
each other

Instead of implementing this functionality into your appn,
a second container can be started that has this
functionality implemented.

* Software to manage network traffic → called proxy.



* Istio Architecture:

When a service mesh is used, appn's don't talk to each other directly, but the traffic is routed through the proxies instead.

- Popular service meshes → istio and linkerd
- Proxies in a service mesh form the data plane.
 - ↓
 - networking rules are implemented and shape the traffic flow.
- Networking rules are managed centrally in the control plane of the service mesh.
 - ↓
 - define how traffic flows from one service to another and what config should be applied to the proxies
- Config file: tells → service mesh that services should always communicate encrypted.
 - ↓
 - uploaded to the control plane and distributed to the data plane to enforce the new rule.
- Service Mesh Interface (SMI): define a specⁿ on how service mesh from various providers can be implemented.
 - ↓
 - Standard for Service mesh integⁿ with providers Kubernetes

* Service Mesh:

- ⌚ - is a way to control how different parts of an app share data with one another.
- dedicated infrastructure layers built right into an app
- Each part of an app → called service

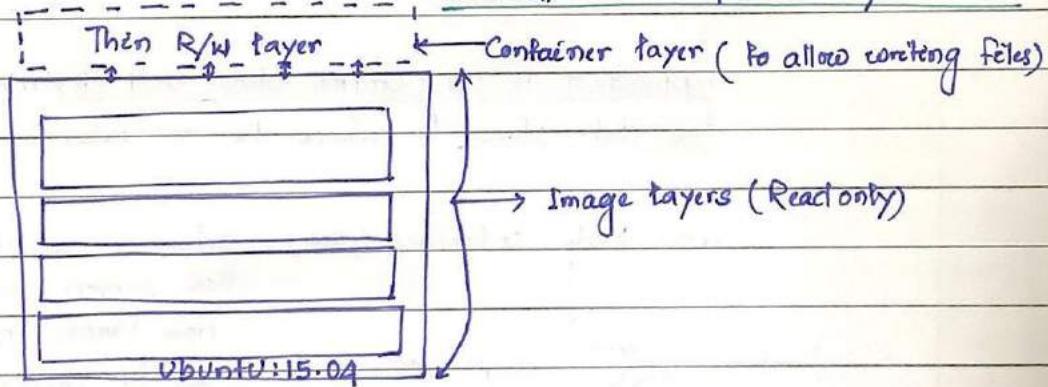
- Service mesh is a way to comprise a large no. of discrete services into a functional appln.
- Designed to handle a high volume of network-based inter-process commn among appln infra. services using APIs.
- Common features provided by service mesh:

Service discovery
load balancing
encryption
failure recovery

* Storage:

- ^{From} Storage perspective: containers have a major disadvantage.

when the container is stopped or deleted, the read-write layer is lost.

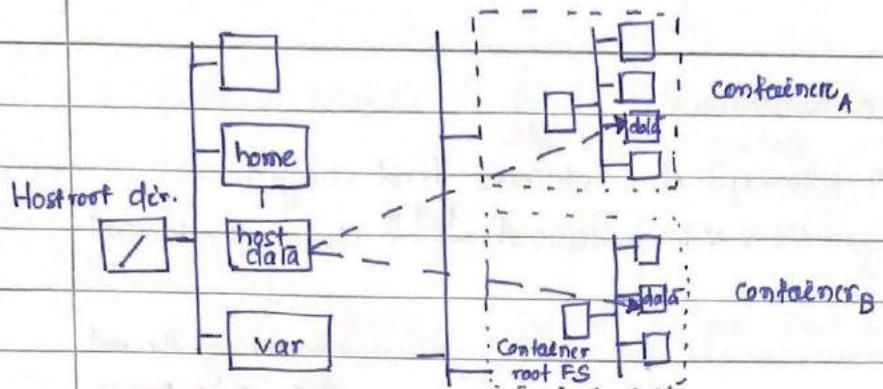


Container
(based on ubuntu:15.04 image)

- R/W layer is lost when the container is stopped or deleted.
- To persist data, you need to write data to disk.
- A container need to write data to a volume on a host.

Directories that reside on a host $\xrightarrow{\text{passed}}$ through into the container file system.

- Data is shared between 2 containers on the same host.



- Challenges:

- a) Persisting the data on the host where the container was started → when you orchestrate a lot of containers
- b) Data needs to be accessed by multiple containers that are started on diff. host systems
- c) When a container gets started on a diff. host, it still should have access to its volume.

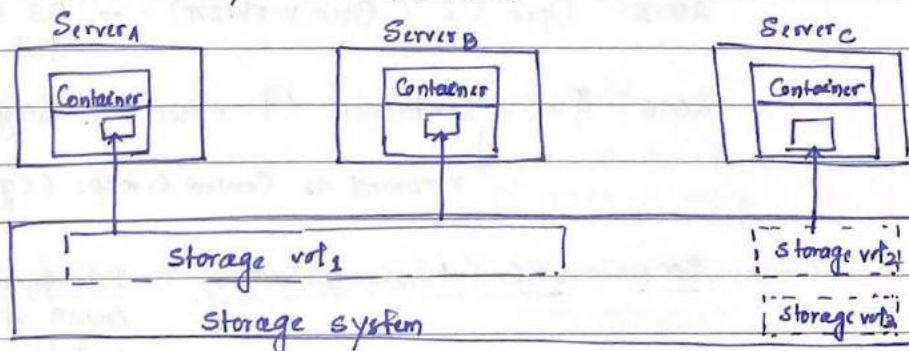
(problem mitigation)

Kubernetes → require → robust storage system that is attached to the host servers.

Container_A and Container_B can share a volume to read and write data.

- Container Storage Interface (CSI):

Standard for exposing arbitrary block and file storage systems to containerized workloads on container orch'n sys. like Kubernetes



- 1. Host based persistence
- 2. Implicit per container storage
- 3. Explicit shared storage

- History of Containers:

1979: concept of containers first emerged.

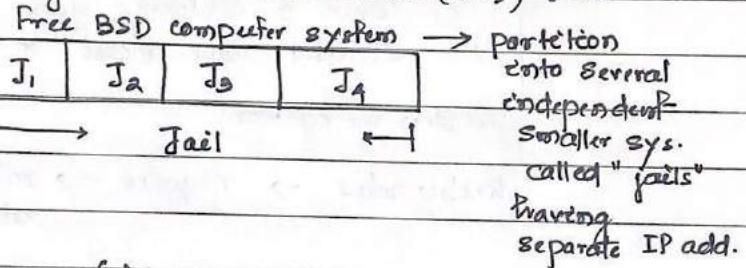
Unix V7 development led to chroot system call



changing the root
dir. of a process
and its children
to a new loc in the FS.

Beginning process isolation → segregating
file access for
each process

2000: Free Berkeley software distribution (BSD) Jails



2001: Linux Vserver (Linux Virtual Server)

↓
jail mechanism that can partition
resources on a comp. sys.

2004: Solaris Container (combines sys. resource controls
and boundary separation
provided by zones)

2005: Open VZ (OpenVZ) → OS level virtualization
tech. for Linux

2006: Process Containers (launched by Google)

↓
renamed as Control Groups (cgroups)

2008: LXC (Linux Containers) → 1st most complete
impl of Linux
container manager

2011: Warden (Cloud Foundry started Warden which can isolate envs on any OS running as a daemon and providing an API for container mgmt.)

2013: LMCTFY (Let me Container That For you)

Open source version of Google's container stack, providing Linux app¹⁰ containers.

2013: Docker → used LXC in initial stages and later replaced container mgmt. with its own library, libcontainer.

2016: Container Security realized → DevSecOps

2017: Maturity on Container Tools (Kubernetes — CNCF - 2016)

2018: CSPs offer managed Kubernetes service

2019: New Runtime engines started replacing Docker runtime engine → containerd
CRI-O

- Best Practices for building Containers:

1) Package single app per container
2) Properly handle PID 1, signal handling and zombie process
(init system)

3) Optimize for the Docker build cache

↓
(accelerate the building of container images)

4) Remove unnecessary tools (to protect your apps from attackers)

5) Build the smallest image possible

6) Reduce the amount of clutter in your image

7) Try to create images with common layers

8) Scan images for vulnerabilities

google/cloud-srlik: 9) Properly tag your images

193.0.0 10) Carefully consider whether to use a public image

↓
image name

image tag

- Play with Docker (PWD) → ^{Project} sponsored by Docker

allows users to run Docker commands
in a matter of secs.

→ Free Linux VM in browser → Build and
run containers
and create
clusters in
Docker Swarm Mode

* The usage of containers and VMs are ^{not} mutually exclusive.

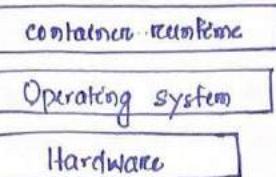
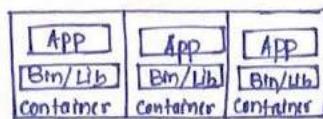
(Two or more
events that
cannot happen
simultaneously)

* Container isolation achieved by namespaces and groups

* Service Discovery → Automatic process of
discovering services on a network

CONTAINER ORCHESTRATION

- Containers: are an app^{ic}entric method to deliver high-performing scalable apps on any infra. of your choice.



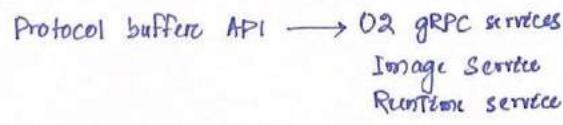
- * encapsulate micro-services and their dependencies but do not run them directly.
- * Containers run container images.

- Container Orch^{ic} Tools:

- * Marathon → framework to run containers on Apache Mesos
- * Nomad → HashiCorp

- Container Runtime Interface (CRI):

- * Container Run Time → at the lowest layers of a K8s cluster, CRT is the S/I that starts and stops containers.
- * CRI: New plug-in API for container runtime in K8s
- * CRI: consists of protocol buffers and gRPC API and libraries.
- * Kubelet - - - - - ^{Communication} CRI Shim over UNIX Sockets using the gRPC framework.
(gRPC Client) (Container runtime)
(gRPC server)



- Container Network Interface (CNI):

- * CNI (a CNCF proj) → specif^{ic} + libraries

- Container Storage Interface (CSI):

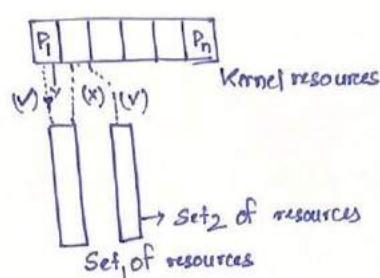
- * Using CSI, 3rd party storage providers can write and deploy plugins exposing new storage sys. in K8s. without ever having to touch the core Kubernetes code.
- * CSI Adoption by container orchestrators:
 - Cloud Foundry
 - K8s
 - Mesos
 - Nomad
- * Once you have deployed a CSI to a K8s cluster, it is available for use with persistent volumes (PV), storage classes and persistent volume claims (PVCs).

Service Mesh Interface (SMI):

- * SMI → specⁿ that covers the service mesh capabilities.
 - a) traffic policy
 - b) traffic telemetry
 - c) traffic mgmt.
- * defines a standard set of K8s custom resource definitions (CRDs) and APIs for service meshes.
- * open source specⁿ for interoperable service meshes on K8s

- Namespace:

- * Linux network namespaces are a Kernel feature allowing us to isolate network envs through virtualization.
e.g. Using network namespaces, you can create separate network interfaces and routing tables that are isolated from the rest of the system and operate independently.



It is a feature of the Linux Kernel that partitions Kernel resources such that one set of processes sees one set of resources while another set of processes sees a different set of resources.

- Cgroups:

- * provides resource limiting, prioritization, accounting, control

- runc:

- * It's a process level tool and not designed with an end user in mind.
- * Has the ability to run containers without root privileges.
- * Command line client for running apps packaged according to the OCI format.
- * It's a Container runtime originally developed by Docker and later extracted out as a separate open source tool.

- Dockerfile:

- * Text document that contains all the commands a user could call on the command line to assemble an image.
- * docker build command builds an image from a Dockerfile and a context.
The build's context is a set of files at a specified location or URL.
- * FROM command → sets the base image for the rest of the instructions.
MAINTAINER command → author of the generated images

RUN Command → used to execute any commands

* Alpine images tightly controlled and small in size → good for Go
bad for Python

* Privileged container in Docker is a bad idea.

↓

→ Containers that have all of the root capabilities of a host machine

→ One use case of a privileged container : Running a docker daemon inside a docker container.

* Do not expose unused ports

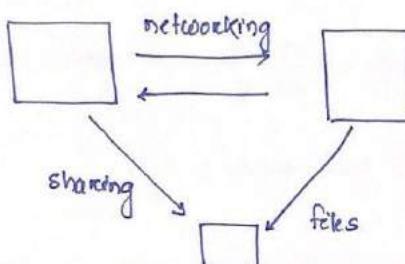
* Do not run SSL within containers

* do not share the host's network namespace

Best practices

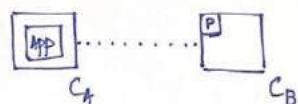
- Docker runtime security

- Networking:



* Two containers can talk to each other in one of the 2 ways.

- Communication through networking → containers can send and receive requests to other app's using networking.
- Sharing files on disk



An app running in one container will create a network container connection to a port on another container.

* Docker comes with network drivers geared towards different use cases. The common network types are:

- Bridge - limited to containers within a single host running the Docker engine. port mapping needs to be configured. difficult Docker network type.
- Overlay - multi-host network comm, uses network tunnels to communicate between hosts
- macvlan - connect Docker containers directly to the host network interfaces through Layer 2 segmentation.

- * Network Plugins in K8s come in a few flavors:
 - a) CNI Plugins
 - b) Kubenet plugin - default network plugin
- * Flannel - overlay network provider that can be used with K8s
- * Cilium - most popular open-source networking and networking security solution for K8s.

DNS Service:

- * DNS is a service that translates domain names into IP addresses.
- * DNS has been designed to use port 53.
- * When you deploy a container on your network, if it cannot find a DNS server defined in `/etc/resolv.conf`, by default it will take on the DNS configured for the host machine.
- * DNS Server takes a

FQDN (www.example.com) $\xrightarrow[\text{IP add.}]{\text{resolves}}$ 93.184.216.34

Service Mesh:

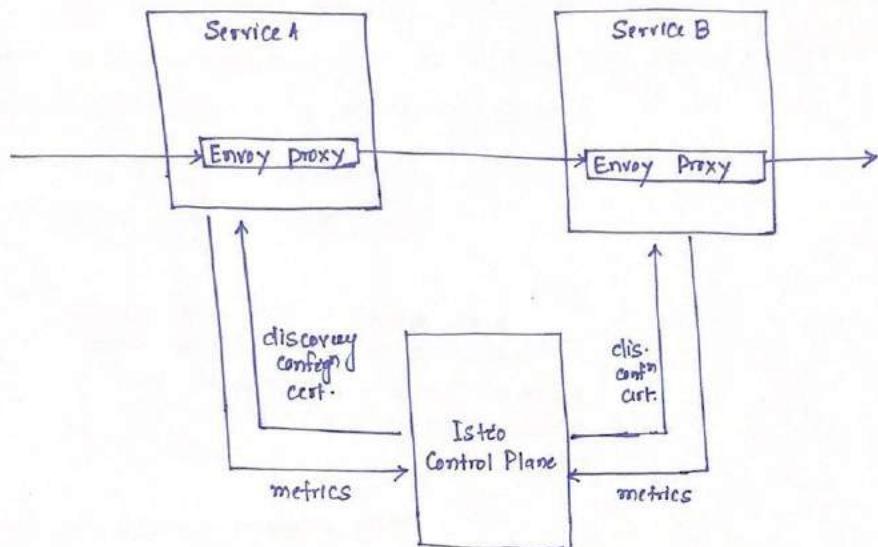
- * A Service mesh is a way to control how diff. parts of an app share data with one another.
- * A Service mesh is built into an app as an array of network proxies.
- * Individual proxies that make up a service mesh are called "sidecars", since they run alongside each service, rather than within them.
- * SM: manages all service-to-service comm within a distributed S/I sys. It accomplishes via the use of "sidecar" proxies that are deployed alongside each service through which all traffic is transparently routed. Sidecars enable service req. to flow through the app.
- * Popular service meshes: Istio
Linkerd
Consul
Kuma
Mesher

Supporting proxies: Envoy
HAProxy
NGINX
MOSN

- * SM: manages "east-west" traffic within a microservices based S/I system.

network traffic
within a datacenter,
network or K8s cluster

- * SM works with a service discovery protocol to detect services as they come up.
- * Istio backed by Google, IBM and Lyft is the best-known service mesh arch. K8s is the only container orchⁿ framework supported by Istio.
- * Proxies → data plane
management processes → control plane
- * Istio Mesh:



- SM: dedicated infra. layer that you can add to your applⁿs.
- SM: addresses more complex operational requirements like
 - A/B testing
 - canary deployments
 - access control
 - encryption
 - end-to-end authentication
- Control Plane:
 - Secure service-to-service commⁿ in a cluster with TLS encryptⁿ
 - Automatic load balancing
 - (Features)

takes your desired config and its view of the services
- Data Plane: communication between services
SM uses a proxy to intercept all your network traffic, allowing a broad set of applⁿ aware features based on config you set.

- Istio extends K8s to establish a programmable app-aware network using the Envoy proxy.
- Istio is platform independent and is designed to run on
 - Cloud
 - On-Premises
 - K8s
 - Mesos
- Istio: Control plane
- Envoy: data plane agent

Service Discovery:

- * Process of automatically detecting devices and services on a network
- * Service discovery protocol (SDP) is a networking standard that accomplishes detection of networks by identifying resources.
- * 02 types of service discovery:
 - i) Server-side
 - ii) client-side
- * Service discovery implementations:
 - i) DNS (NGINX proxy which can periodically poll for service discovery)
 - ii) Key-value store and side-car (Consul or Zookeeper)
 - iii) Specialized service discovery and library/sidecar (Eureka)

Working with Kubernetes

Kubernetes Fundamentals

- Pod: smallest compute unit on Kubernetes

- Kubernetes: config mgmt.
cross-node networking
routing of external traffic
Pod balancing
scaling of the pods

- Kubernetes Objects:

★ * Kubernetes objects are persistent entities in the Kubernetes system.

(represent the state of the cluster)

describe

- what containerized apps are running (and on which nodes)
- the resources available to those apps
- the policies around how those apps behave, such as
 - restart policies
 - upgrades
 - fault tolerance

★ * Kubernetes obj. → record of entity → once you create that obj, Kubernetes system will constantly work to ensure that the obj. exists.

* To work with Kubernetes objects -- whether to create, modify or delete them

→ REST

Kubernetes API (core of Kubernetes' control plane)

★ * Kubectl CLI: makes the API call for you

* In Kubernetes, deployment obj. → represent an app running on your cluster

- * Every Kubernetes obj. includes 02 nested obj. fields that govern the obj.'s config.

i) obj. spec ii) obj. status

you have to
set this when
you create the obj.
providing a descn
of the char. you
want the resource
to have; its desired
state.

describes the current
state of the obj.

- * When you create an object in Kubernetes, you must provide

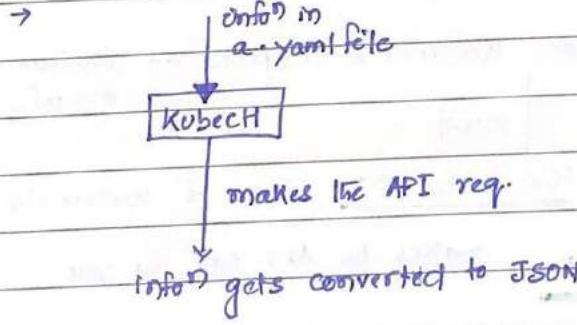
a) obj. spec → desired state

b) Basic info. about the obj. e.g. name

→ When you use the Kubernetes API to create the obj.,
(either directly or via Kubectl),

API request must include as JSON in the req. body
that info.

(JavaScript Obj. notn,
standard text based
format for
representing
structured data
based on Javascript
obj. syntax.)



→ Sending some data
from server to client
so that it can be
displayed on a
web page)

→ which version of the Kubernetes API you are using to create this obj.

apiVersion: apps/v1

kind: Deployment

what kind of obj. you want to create

metadata: → data that helps uniquely identify the obj.

name: nginx-deployment

spec: → what state you desire for the obj.

Selector:

app: nginx

replicas: 2 # tells deployment to run 2 pods matching the template

template:

app: nginx

spec:

Containers:

- name: nginx
- image: nginx:1.19

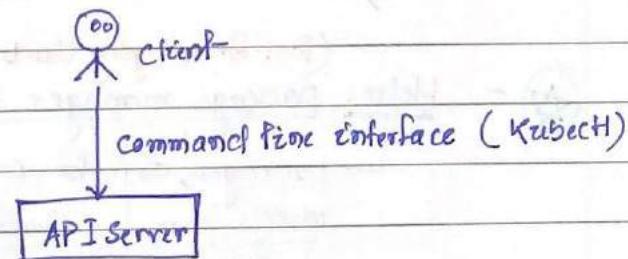
ports:

- containerPort: 80

→ Create a deployment: kubectl apply -f <https://k8s.io/examples/application/deployment.yaml> --record

o/p: deployment.apps/nginx-deployment created

- Interacting with Kubernetes:



① - Kubectl api-resources

name	shortnames	api-version	namespace	Kind
configmaps	cm	v1	true	configmap
namespaces	ns	v1	false	namespace
Pods	po	v1	true	pod

- Kubectl explain pod → get docⁿ of various resources

Kind : Pod

Version: v1

[Pod → colleⁿ of containers that can run on a host. This resource is created by clients and scheduled onto hosts.]

- Kubectl explain pod-spec

② - Kubectl --help (Kubectl → Kubernetes ~~cluster manager~~)

↓ overview of all the commands

③ - To create an object in Kubernetes, from a YAML file.

↓

Kubectl create -f <yourfile>.yaml

- Tools for interaction with K8s:

- 1) Kubernetes dashboard
- 2) cldrailed/K9s
- 3) Lens
- 4) VMWare Tanzu Octant

(Tool for managing charts, tool that streamlines installing and managing K8s apps)

④ - Helm: package manager for Kubernetes

↓ packages objects in so called charts which can be shared with others via a registry.

(Helm charts is a collection of files that describe a related set of K8s resources)

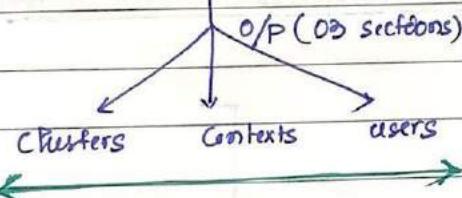
Helm is the K8s equivalent of yum or apt.

- * Open-Source proj.
- Artifact Hub: web-based app that enables finding, installing, and publishing K8s packages and configs for CNCF proj.

- Kubectl: command-line tool to control K8s clusters.

For config, Kubectl looks for a file named config in the \$HOME/.kube directory.

\$kubectl config view → current Kubectl config



1) define a K8s object in a YAML file (Pod running a nginx container on port 80)

\$ vim pod.yaml. [Object spec]

2) \$ kubectl create -f pod.yaml → Pod/nginx created

3) \$ Kubectl get pods

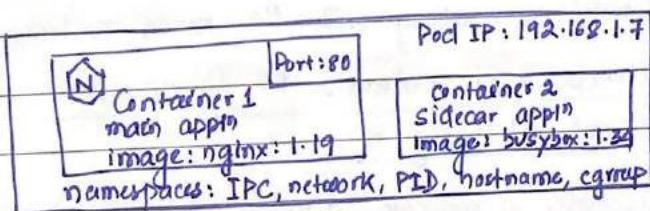
name ready status restarts age
nginx 1/1 running 0 10s

4) \$ kubectl delete pod nginx → Pod "nginx" deleted

5) \$ kubectl get pod → no resources found in default namespace.

6) \$ kubectl explain --help

Pod Concept:



- * A pod is a group of one or more containers with shared storage and network resources, and a specifn of how to run the containers.
- * A pod's contents are always co-located and co-scheduled and run in a shared context.

- * Pod describes a unit of one or more containers that share an isolation layer of namespaces and groups.

- * smallest deployable unit in Kubernetes \Rightarrow Kubernetes is not interacting with containers directly.

- * all containers inside a pod share an IP address and can share via the filesystem.

- * multiple containers share namespaces to form a pod.

* apiVersion: v1

kind: pod

metadata:

name: nginx-with-sidecar

spec:

Containers:

- name: nginx

image: nginx:1.19

Ports:

- ContainerPort: 80

- name: Count

image: busybox:1.34

simple pod
object with
two containers

- * Sidecar Container:

\rightarrow a utility container in a pod that's loosely coupled to the main application container

\rightarrow shared runs along with the main container in the pod.

\rightarrow extends and enhances the functionality of current containers without changing it

\rightarrow Resembles a sidecar attached to a motorcycle

- Init Container:

- * it's the one that starts and executes before other containers in the same pod.
- * meant to perform initialization logic for the main app hosted on the pod.
- * Always run to completion

- Create a pod: `$ kubectl apply -f https://k8s.io/examples/pods/simple-pod.yaml`

- One CPU in Kubernetes \approx 1 vCPU/core for cloud providers and 1 hyperthread on bare-metal Intel processors

 - **liveness probe:** Configure a health check that periodically checks if your app is still alive

- **Security Context:** Set user & group settings, as well as kernel capabilities.

- `$ docker run --detach nginx:1.19` \rightarrow run a container

`$ docker ps` \rightarrow check the running status of the container on that host
Container ID Image Command Created Status Ports Names
nginx:1.19

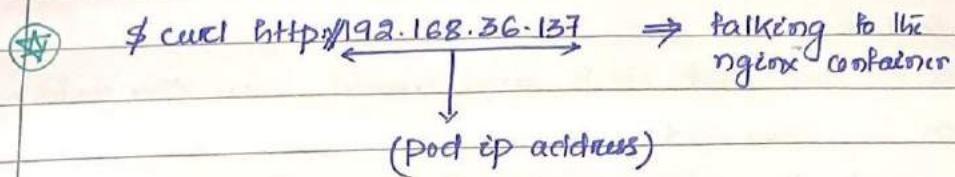
`$ kubectl run nginx --image=nginx:1.19` \rightarrow run a pod

\rightarrow Pod/nginx created

`$ kubectl get pods`

\downarrow
Name Ready Status Restart Age
nginx 1/1 running 0 5s

`$ kubectl describe pod pod nginx` \rightarrow more details about nginx pod.



↙ \$ kubectl create -f pod-two-containers.yaml

↙ pod/nginx-with-sidecar created

↙ \$ kubectl get pods

↙ name ready status restarts age

nginx	1/1	running
nginx-with-sidecar	2/2	running

↙ 2 out of 2
 containers are
 running

Workload Objects:

* ReplicaSet: a controller obj. that ensures a desired
 ↗ number of pods is running at any given time.

↙ ReplicaSets can be used to scale out apps and
 improve their availability. They do this by
 starting multiple copies of a pod def.

ensures how many replica of pod should be running.

↗ To manage the scaling of pods, Kubernetes uses an
 API object called a ReplicaSet.

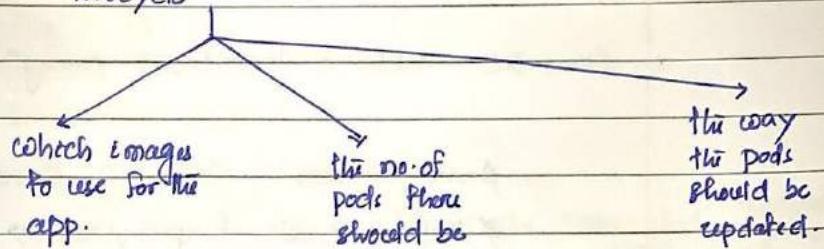
* Deployment: most feature-rich object in Kubernetes.

↗ describe the complete app's life-cycle

↙ they do this by managing
 multiple ReplicaSets that get
 updated when the app is changed
 by a new container image.

→ Resource obj. in Kubernetes that provides declarative updates to apps. (for pods and replica sets)

→ Deployment - allows you to describe an app's lifecycle



→ Deployments are perfect to run stateless apps

(app. prog. that doesn't save client data generated in one session for use in next session with that client)

* Stateful Set: - work load API object used to manage stateful apps.



- Controller that helps you deploy and scale groups of Kubernetes pods.

(program that saves client data from the activities of one session for use in the next session)



- try to retain IP addresses of pods and give them a stable name, persistent storage and more graceful handling of scaling and update.

e.g. databases
(mongodb
elasticsearch)
etc.



* daemonSet: → ensures that a copy of the pod runs on all (or some) nodes of your cluster

→ perfect to run infra.-related workload e.g.

kind=daemon-set

monitoring or logging tools

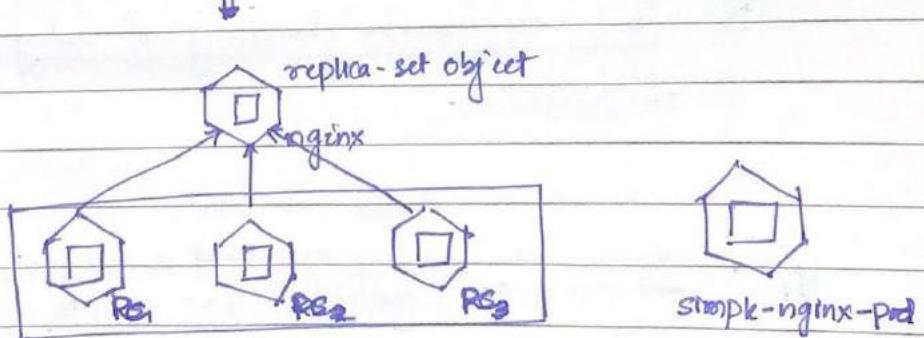
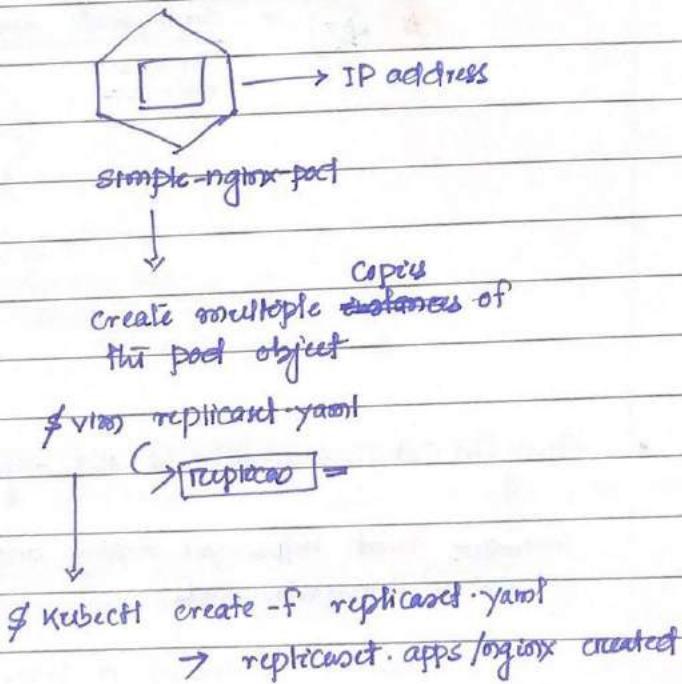
- * Job: creates one or more pods that execute a task and terminates afterwards.

Job objects → one-shot scripts → database migrations or administrative tasks

- * Cron Job: adds a time-based config to jobs.

\$ vi `pod.yaml`

\$ kubectl create -f `pod.yaml` → pod/simple-nginx-pod created



\$ kubectl get pods

name	ready	status	restarts	age
------	-------	--------	----------	-----

nginx-PC ₁	1/1			
-----------------------	-----	--	--	--

nginx-PC ₂	1/1			
-----------------------	-----	--	--	--

nginx-PC ₃	1/1			
-----------------------	-----	--	--	--

simple-nginx-pod	1/1			
------------------	-----	--	--	--

\$ kubectl get pods -o wide

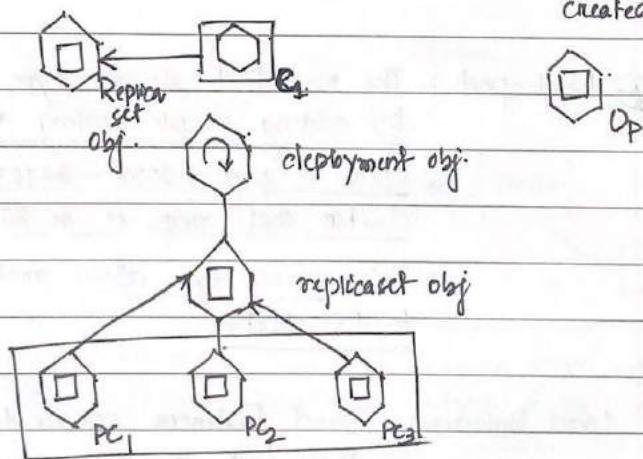
name	ready	status	restarts	IP	node	nominal node	readiness gates
------	-------	--------	----------	----	------	--------------	-----------------

\$ kubectl scale --replicas=10 deployment/nginx → (10 copies of the pod created)

↓
(replicaset.apps/nginx scaled)

\$ vi deployment.yaml (deployment obj manage multiple replica sets)

\$ kubectl create f deployment.yaml ⇒ deployment.apps/nginx created



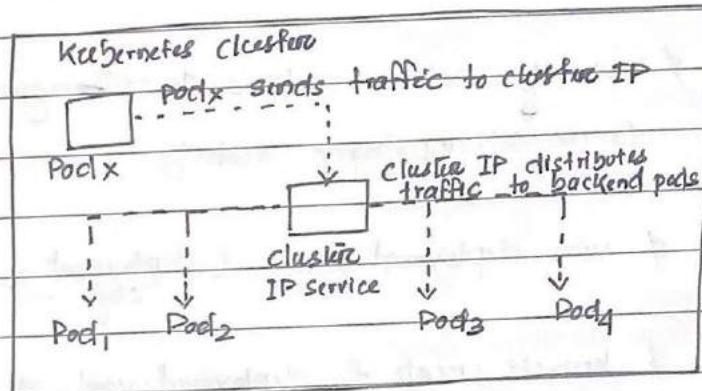
\$ kubectl scale --replicas=10 deployment/nginx

↓
(deployment.apps/nginx scaled)

→ (update the deployment)
\$ kubectl set image deployment/inginx-inginx=nginx:1.20
↓
deployment.apps/inginx image updated

Networking Objects:

Cluster IP: virtual IP inside Kubernetes that can be used as a single endpoint for a set of pods.
↓
default service type. → this service-type can be used as a round-robin load balancer.



Node-Port: The Node-Port service type extends the clusterIP by adding simple routing rules.
→ opens a port (30000 - 32767) on every node in the cluster and maps it to the Cluster IP.
→ this service type allows routing external traffic to the cluster.

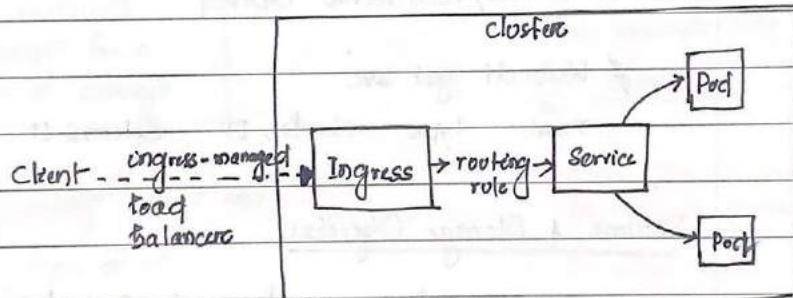
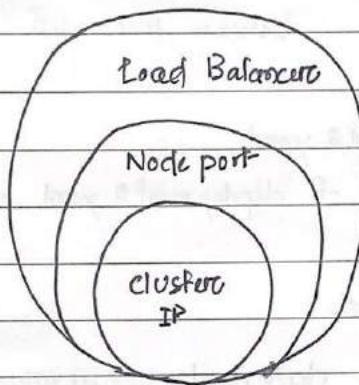
Load Balancers: Load Balancer service type extends the Node-Port by deploying an external Load Balancer instance.
→ will only work if you are in an env. that has an API to configure a Load Balancer instance like GCP, AWS, Azure

External Name: a special service type that has no routing

→ Ext. Name uses → create a DNS alias
 K8s internal
 DNS service

used to resolve
rather a complicated
host name

e.g.
my-cool-database-az1-vcf123.
cloud-provider-e-eke.com



* Ingress obj. provides, a means to expose HTTP and HTTPS routes from outside of the cluster for a service within the cluster.

↓
configuring routing rules
that a user can set and
implement with an
Ingress controller.

Features [load balancing, TLS offloading/termination,
 name-based virtual hosting, path-based routing]

K8s Provele, a cluster internal firewall with the Network Policy concept.

↓
simple IP firewall (OSI layer 3 or 4)
that can control traffic based on rules.

e.g. A typical use-case for Network Policies would be restricting the traffic between two diff. namespaces.

\$ vim deployment2.yaml

\$ kubectl create -f deployment2.yaml → deployment.apps /echoserver created

↓
03 pods

\$ kubectl expose deployment echoserver --port=8080

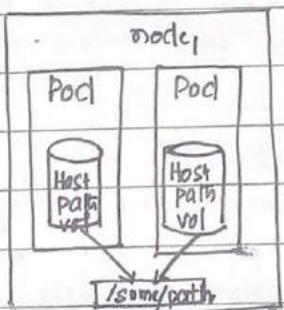
↓
service/echo-service exposed

\$ kubectl get svc

name	type	cluster-IP	external-IP	port(s)	age
------	------	------------	-------------	---------	-----

- Volume & Storage Objects:

K8s → volume mount → as a part of a pod



sharing data between multiple containers within the same pod.

Prevents data loss when a pod crashes and is restarted on the same node.

Persistent storage:

Cloud block storage - Amazon EBS

Google Persistent Disks

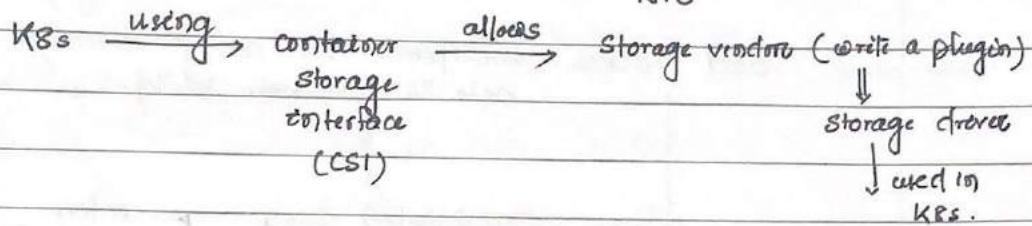
Azure disk storage

Storage systems -

Ceph

GiggleFS

NFS



02 storage objects

- * PVCs consume PV resources.
- * PVCs are requests for PVs.

* Provide APIs for users & admins to manage and consume persistent storage

* Provisioned by admin. or dynamically provisioned by storage classes

* Volume plugins like volumes

* resources in the cluster

Persistent volumes (PV)

an abstract description for a slice of storage

obj. config holds info.

type of vol.

vol. size

access mode and unique identifiers

Persistent Volume Chain (PVC)

If the cluster has multiple PVs,

user creates PVC

* PVC can be used to request storage.

* 03 access modes

ReadWriteOnce

ReadOnlyMany

ReadWriteMany



Rook: Open-source cloud-native storage orchestrator for K8s.

Providing the platform,

framework, and

support for a diverse set of storage softs

automates the task of a storage admin.

- Configuration Objects:

App.

config files, conn^o to other services,
DB sys, storage

config^o

(conn^o strings)

Bad practice: Incorporate the config directly
into the container build

any config change → entire
image need
to be
rebuilt
And the entire
container or pod
to be redeployed.

solved

K8s → decoupling the config^o
from the pods with a config-map.

API
object

stores
config-files
or
variables
as key-value pairs.

* Pods can consume config-maps as
env. variables, command line args or as
config^o files in a volume.

* config-maps allow you to
decouple env. specific info from
your container images.

* data stored in a CM cannot
exceed 1 MiB.

* Usage: configure settings for
containers running in a pod
in the same namespace.

use in 2 ways

mount
a config-map
as a
volume
in pod

map
variables
from a
config-map
to
env. variables
of a pod

* Opaque Secret: arbitrary user-defined data
default Secret type

* Kustomize provides resource Generators to create secrets and configurations.
→ you don't need to include confidential data in your app's code.

- Secrets: K8s objects to store sensitive info like creds, token, or a key

↳ base64 encoded

cloud-native env. Secret mgmt. tools have emerged.

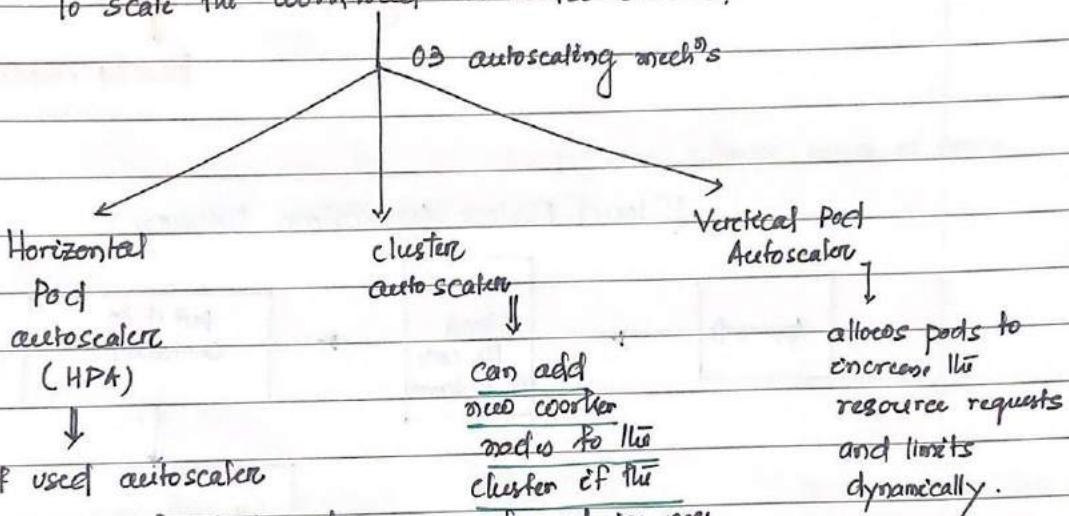


(HashiCorp Vault)

* Stored unencrypted in etcd. * Bootstrap token secrets → Automate node regis.

Auto Scaling Objects:

To scale the workload in a K8s cluster.



installing of an add-on server

Called metrics Server.

replace metrics Server with Prometheus adapter for Kubernetes Metrics APIs.

use custom metrics

* Coregularities in the cluster pods

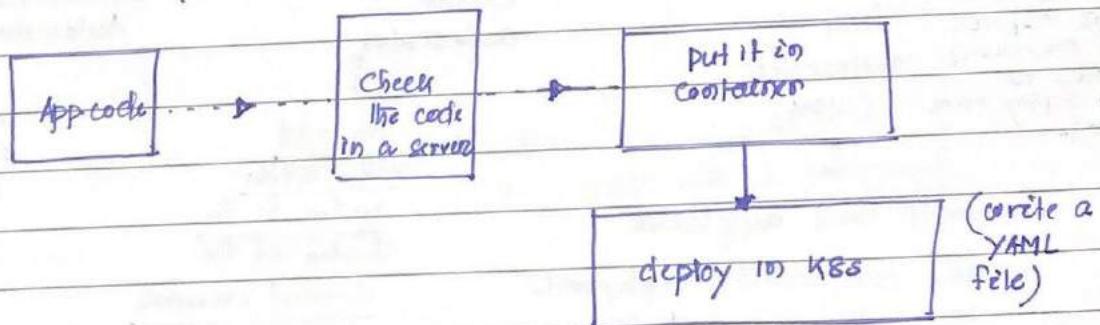
KEDA: (Kubernetes-based Event Driven Autoscaling)

proj. KEDA can be used to scale the k8s workload based on events triggered by external systems.

2019 (Microsoft - Red Hat partnership)

- * K8s objects can be described in a data serialization lang. called YAML
- * Main difference between ConfigMaps and Secrets
 - ↓
Base64 encoding

(Cloud-Native Application Delivery)



Application Delivery Fundamentals:

- * Git - standard version control system

↳ cheap local branching
staging areas
multiple workflows

- * IaC (Infra-as code) → process of managing and provisioning data centers through machine readable def'n files. rather than phy. H/W config.

Ansible Autom8 Platform

Playbooks are used

to describe the desired state of your infra. → config/network policies/security

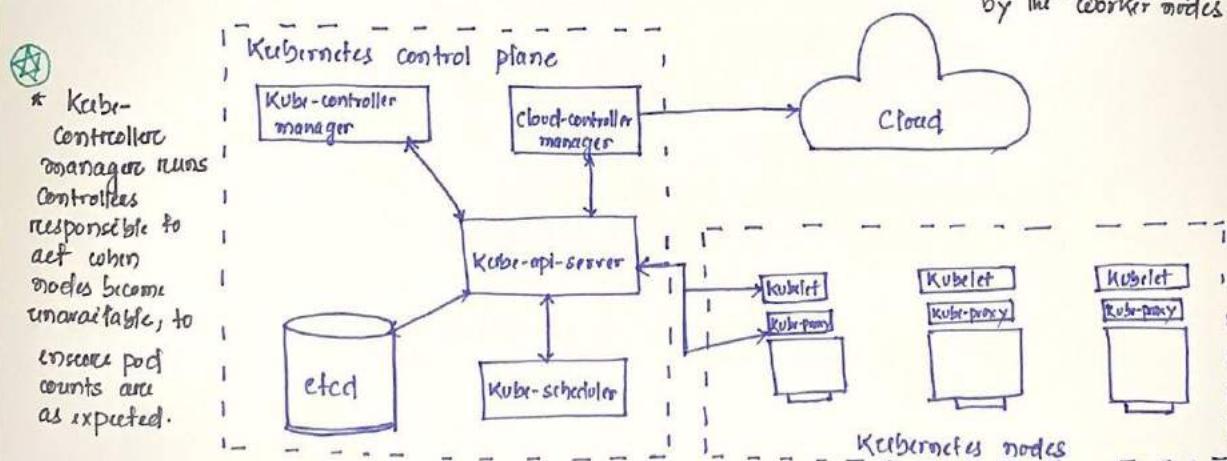
↑ (can be described as code)

KUBERNETES FUNDAMENTALS

- Kubernetes: container orch¹⁰ platform
 - ↳ automate deployments
 - ↳ Scaling and mgmt. of containerized workloads
 - ↳ written in Go language
- Kubernetes: developed by Google in 2013
 - ↳ highly inspired by Google Borg System
 - ↳ transitioned to CNCF in 2014 (1st proj.)
 - ↳ current version of Kubernetes - 1.23
- Architecture: Kubernetes clusters consist of two different server nodes.
 - 1) Control Plane (nodes)
 - 2) Worker nodes

(manage the cluster and the worker nodes) * brain of open

 - * One or more master nodes as part of the control plane
 - * Contain various components which manage the cluster and control various tasks like deployment, scheduling and self-healing of Containerized workloads.
 - * apps run in your cluster.
 - * only job of worker nodes (running of apps)
 - * Network traffic between client users and the containerized apps deployed in pods is handled directly by the worker nodes.



- * Control-plane nodes typically host the following services:
 - 1) Kube-apiserver: Centerpiece of Kubernetes
 - all other components interact with the api server and it is where users should access the cluster.

2) etcd : database which holds the state of the cluster.

↳ standalone proj. and not an official part of Kubernetes.

3) Kube-scheduler : When a new workload should be scheduled, the Kube-scheduler chooses a worker node that could fit based on CPU and memory.

4) Kube-controller-manager : contains many non-terminating control loops that manage the state of the cluster.

5) Cloud-controller-manager : can be used to interact with the API of the CSPs to create external resources like

load balancers
Storage or
Security groups

* Components of worker nodes :

1) Container runtime : responsible for running containers on the worker node.

e.g. Docker → popular choice
↑

replaced by containerd

2) Kubelet : a small agent that runs on every worker node in the cluster.

→ receives pod def's from the API server, and interacts with the container runtime on the node to run containers associated with the Pod.

→ monitors the health and resources of pods running containers.

3) Kube-proxy : a network proxy that handles inside and outside commⁿ of your cluster.

→ relies on the networking capabilities of the underlying OS for managing traffic flow.

→ responsible for dynamic updates and maintenance of all networking rules on the node.

* Control Plane :

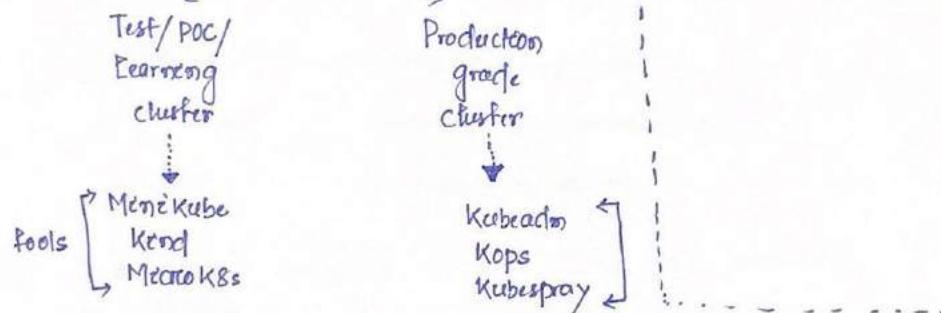
1) If not available → apps that are already started on a worker node will continue to run.

→ scaling, scheduling new apps etc: (not possible) (2)

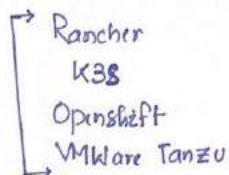
- * NS provides a mech. for isolating groups of resources within a single cluster.
 - * NS → way to divide cluster resources between multiple users
 - * Kubernetes namespace: can be used to divide a cluster into multiple vertical clusters, which can be used for multi-tenancy when multiple teams share a cluster.
- Viewed like a directory on a comp. where you can organize objects and manage which user has access to which folder.

- Kubernetes Setup: - - - - -

- * Setting up a Kubernetes cluster



- * Vendors started packaging Kubernetes into a dist^{ro} and offer commercial support.



- * CSPs:
 - Amazon: (EKS) → Elastic Kubernetes Service
 - Google: (GKE) → Google Kubernetes Engine
 - Microsoft: (AKS) → Azure Kubernetes Service
 - Digital Ocean: (DOKS) → Digital Ocean's managed Kubernetes service

- Demo: Create a node Kubernetes cluster → Kubeadm

Pre-requisites:

- 1) Linux Host
- 2) 2 GB or more of RAM per machine

3) 2 CPUs or more

4) Feel network connectivity between all machines in the cluster

* Installation of Container Runtime:

→ To run containers in Pods, Kubernetes uses a container runtime.

→ By default, Kubernetes uses a container runtime interface (CRI) to interface with your chosen runtime.

→ If you don't specify a runtime,

KubeADM automatically tries to detect an ~~installed~~ container runtime by scanning through a list of well known Unix domain sockets.

Runtime	Path to Unix domain socket
Docker	/var/run/docker.sock
Containerd	/run/containerd/containerd.sock
CRI-O	/var/run/crō/crō.sock

→ Docker takes precedence if both Docker and containerd are detected.

* Installation of packages in all the machines.

★ → KubeADM: command to bootstrap the cluster (getting the initial cluster up and running)

→ Kubelet: primary node agent that runs on each node (starts pods and containers)

→ Kubectl: Kubernetes command line tool → allows you to run commands against Kubernetes clusters.

* Initialize control plane node Sudo
KubeADM init (root privileges)

* Installation of a pod-network: install Calico networking and network policy for on-premises deployments

↓
Install a overlay network

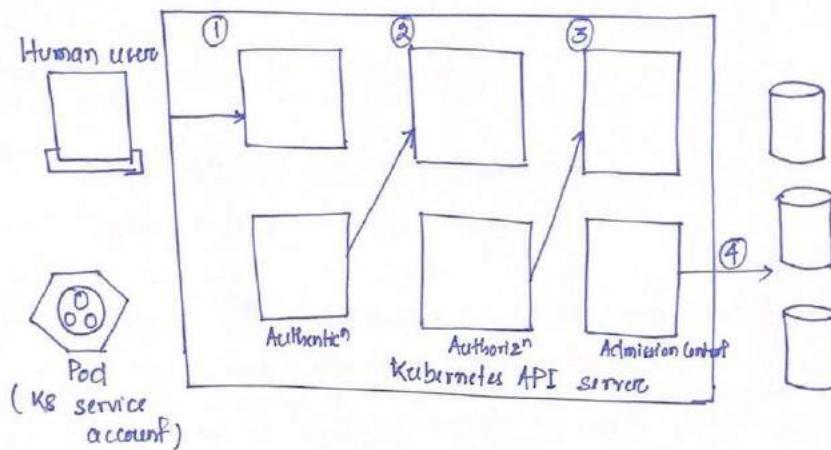
- * `sudo kubeadm join` Command → join a node to the cluster

- * `kubectl get nodes`

name	status	roles	age	version
CP
worker node

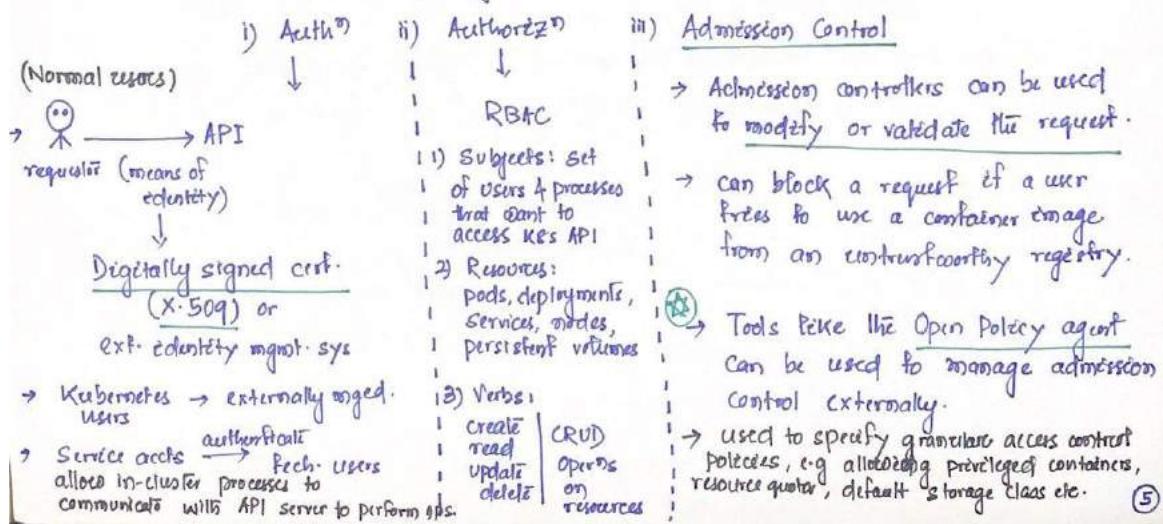
- * `kubectl cluster-info` : info about cluster

- Kubernetes API:



- * Connecting with the cluster → Kubernetes API server

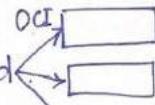
- * Before a req. is processed by the Kubernetes, it has to go through 03 stages:



- * Kubernetes API is implemented as a RESTful interface that is exposed over HTTPS.
 - ↓
 - (Representational State Transfer)
 - ↓
 - Architectural style for an API that uses HTTP requests to access and test data.
 - also known as REST API.
 - most common method for connecting components in micro-service architectures.

Running Containers on Kubernetes:

- * define Pods → smallest compute unit
 - translate Pods → into running containers.
- * Kubelet $\xleftarrow{\text{CRI}}$ containerd $\xleftarrow{\text{}}$ Docker
 - Docker uses a shim app¹⁹ which provides a clear abstraction layer between Kubelet and the container runtime.
 - containerd → lightweight and performant impl²⁰ to run containers.
 - ↓
 - most popular container runtime



CRI-O → created by RedHat with a similar codebase as podman and buildah

Docker → usage of Docker as the runtime for Kubernetes has deprecated, and will be removed in Kubernetes 1.25.

- * Kubelet connects to container runtimes through a plug-in interface (CRI). CRI $\xrightarrow{\text{uses}}$ protocol buffers, gRPC, libraries etc.
 - * containerd/CRI-O: provides a runtime that only contains the absolutely essentials to run containers.
- * frantic CRI shim → aimed at enabling Kubelet to interact with Kata containers.

* CRI implements O2 services.
 Image Service → image related ops
 Runtime Service → pod & container related ops.

* Containerd / CRI-O $\xrightarrow{\text{integrate}}$ With container runtime sandboxing tools



Solve the security problem that comes with sharing the Kernel between multiple containers.

Tools: 1) gvisor 2) Kata containers

Provides an appvm kernel that sits between the containerized process and the host kernel.

A secure runtime that provides a lightweight VM, but behaves like a container.

- Networking: * A container runtime creates an isolated network space for each container it starts.

* 1) Container-to-container comm¹¹ \rightarrow pod concept [On Linux, isolated network space \rightarrow network namespace]

→ When a pod is started, a special proxy container is centralized by the container runtime to create a namespace for the pod.

2) Pod-to-pod comm¹¹ \rightarrow overlay network

3) Pod-to-service comm¹¹ \rightarrow Kube-proxy and packet filter on the node

4) External-to-service comm¹¹ \rightarrow (- do -)

↳ K8s network model treats pods as VMs on a network, where each VM is equipped with a network interface - thus each pod receiving a unique IP add. This model is called "IP-per-pod" (network proxy that runs on each node in your cluster) and encodes pod-to-pod comm¹¹.

* Implement of networking in Kubernetes

03 comp. requirements

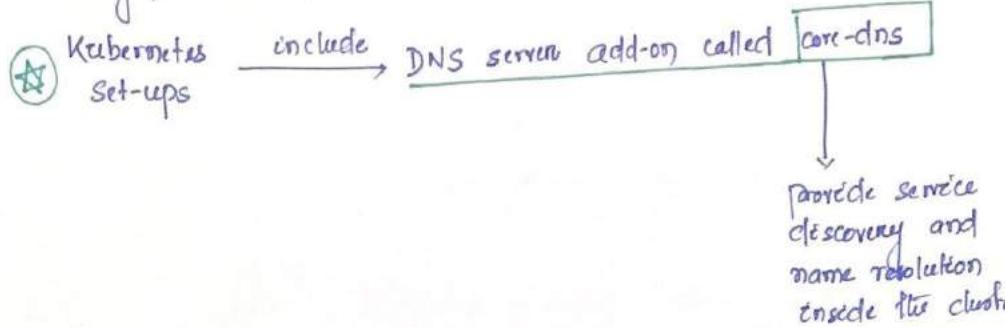
- 1) all pods can communicate with each other across nodes.
- 2) all nodes can communicate with all pods.
- 3) no network address translation (NAT)

* K8s enables ext. accessibility through Services, complex rule defns stored in ipTables on cluster nodes

* To complement networking, you can choose from a variety of and implemented by network vendors:

- 1) Project Calico
- 2) Weave
- 3) Cilium

- * Every pod \rightarrow own IP address \rightarrow no manual config. involved.



Scheduling:

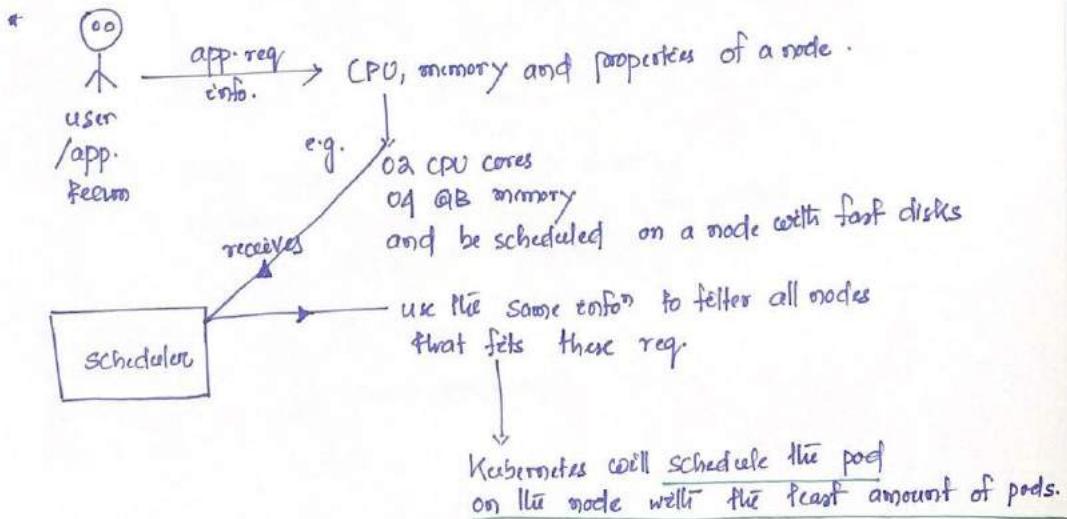
* Container orchⁿ \rightarrow sub-category (Scheduling) \rightarrow process of automatically choosing the right (worker) node to run a containerized workload.

* Kubernetes: Kube-scheduler \rightarrow scheduling

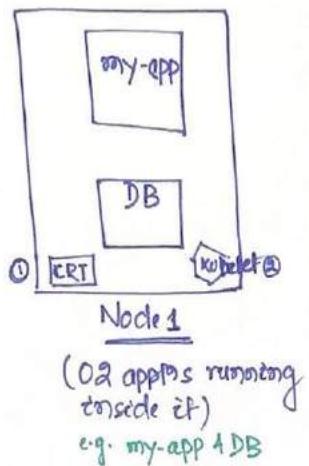
* scheduling process $\xrightarrow{\text{starts}}$ when a new pod is created.

* Kubernetes $\xrightarrow{\text{uses}}$ declarative approach

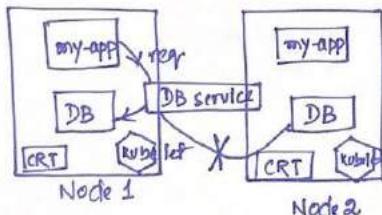
- ★
- pod is only described first.
 - scheduler selects a node where pod ^{actually} will get started by the kubelet and the container runs there.



Kubernetes Architecture



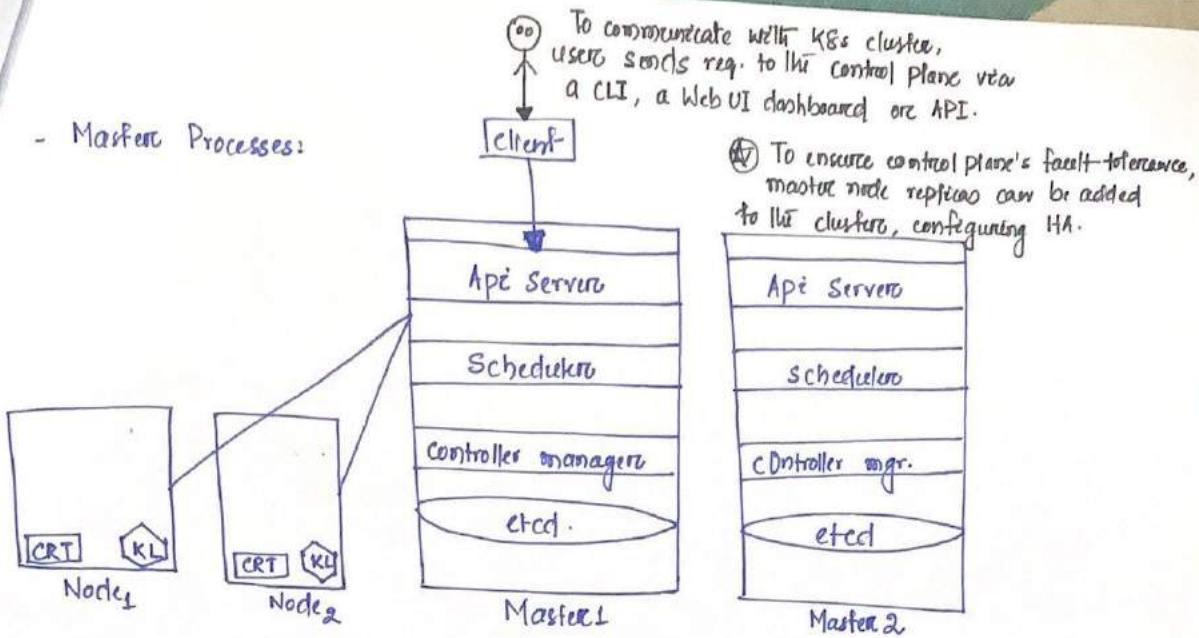
- Each node has multiple pods on it.
 - 03 processes must be installed on each node
 - Worker nodes — computation.
 - * If it's a container engine, basically a S/W component that can run containers on host OS.
 - ! Container runtime
 - 2) Kubelet (node agent)
 - 3) Kube-proxy (proxy)
 - * Setting up namespaces and cgroups for containers
 - * Docker, CRI-O, containerd, rkt (hypervisor based container runtime for K8s)
 - * App's pods have containers running inside it.
 - * Container run time needs to be installed on every node.
 - * Kubelet interacts with the container runtime and the node.
- Takes the config and starts the pod with a container inside.
assigning resources from the node to the container.
- (CPU, memory, storage resources)



- ★ → Comm¹ between pods is through services.
- ★ → Kube-proxy: intelligent routing logic inside that encloses Comm² works in a performance way with low overhead.

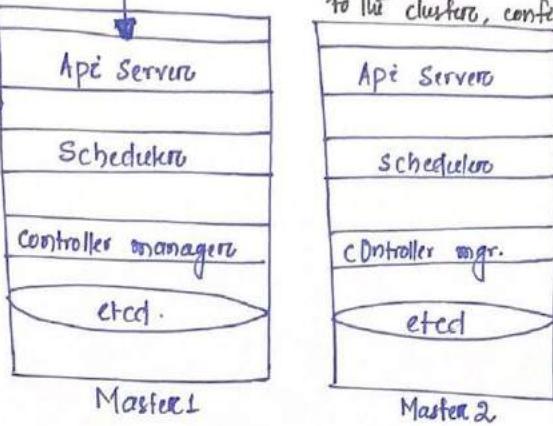
- Schedule pod
 - Monitoring pods
 - Re-schedule/re-start pods
 - join a new node
- Master node (managing processes)

- Master Processes:



To communicate with K8s cluster, user sends req. to the control plane via a CLI, a Web UI dashboard or API.

④ To ensure control plane's fault-tolerance, master node replicas can be added to the cluster, configuring HA.



④ 4 processes run on every master node (run the cluster state and the worker nodes)

- | | | | |
|---|--|---|---|
| 1) Api Server: | 2) Scheduler | 3) controller mgr. | 4) etcd |
| <ul style="list-style-type: none"> wants to deploy a new app in a K8s cluster, you interact with Api Server using some client like UI/K8s dashboard or command line tool like Kubectl. Api server → Cluster gateway, which gets the initial req., may be update to the cluster, or the queries from the cluster. <u>gatekeeper for auth</u> To schedule new pods or deploy new apps, api server receives the req, validates it and feeds it to other processes. Only 01 endpoint to the cluster. Api Server - <u>load balanced</u> Api Server → only master plane component to talk to the etcd store. | <ul style="list-style-type: none"> * intelligent way of scheduling which worker node can handle the running of the pod. * scheduler ⇒ sleep making. * process that does the scheduling of the Kubectl which starts the pod. | <ul style="list-style-type: none"> * detects the state changes of the cluster. * pods die, reschedules the pod and restarts them. * regulates the state of the K8s cluster. * controller's are watch-loops continuously running and comparing the cluster's desired state with the current state. | <ul style="list-style-type: none"> * key-value store of a cluster state. * cluster broken * cluster state → saved. * resources availability cluster health ↓ stored etcd * appn data → not stored in etcd. * etcd ⇒ distributed storage across all master nodes. * Both stacked and external etcd config's support HA config's. * etcd is based on the Raft consensus algorithm. |

- Add new Master/Node Server:

- 1) Get the bare metal server
- 2) install the master/worker nodes processes
- 3) add it to the cluster

- Kubernetes Security Tools:

↑
Popeye
Kube audit
Kube-bench

- * Kubernetes Features:

- 1) Automatic bin packing → automatically schedules containers
- 2) Self-healing → automatically replaces and reschedules containers from failed nodes.
- 3) Horizontal scaling → Based on CPU and memory metrics
- 4) Service discovery and load balancing

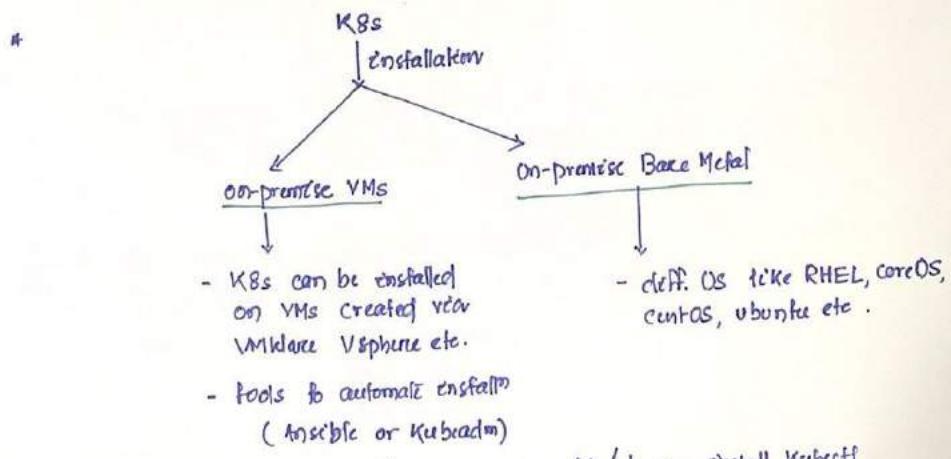
Containers receive their IP address from Kubernetes.

↓
assigns a single DNS name
to a set of containers to add
in load balancing requests
across the containers of the set.

- 5) automated rollouts and rollbacks
- 6) Secrets and config mgmt (K8s manages sensitive data and config details for an app separately from the container image, in order to avoid a re-build of the respective image.)
- 7) Storage orchestration

Kubernetes Configuration:

- * Master installation types are
 - all in one single node (learning, testing)
 - single master and multiple worker (master running a stacked etcd instance)
 - single master with single node etcd, and multiple workers (external etcd instance)
 - multi-master and multi-worker (HA)
 - multi-masters with multi-node and multi-workers]
 - (external etcd instance)
 - (prod env.)
- * Installation tools to deploy single node or multi-node K8s cluster
 - Minikube - easiest and preferred to set up K8s locally.
 - kind
 - Docker Desktop
 - MicroK8s
 - K3s - lightweight K8s cluster from Rancher



- * apt - package manager (native pkg. mgmt) / brew - install Kubernetes on macOS
- * Cloud Environment:

1) Hosted Solutions

- a) Turnkey Cloud Solutions (install K8s on IaaS platform)

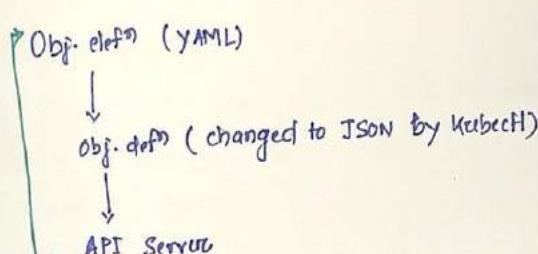
AWS ECA, GCE, Alibaba Cloud

- b) Turnkey On-premise solutions

OpenShift Container Platform

- * Kubeadm: Secure and recommended method to bootstrap a multi-node prodⁿ ready HA K8s cluster
- * Kubespray: (formerly known as Kargo) - can install HA prodⁿ ready clusters on CSPs or Bare Metal.
- * Kops: create, upgrade and maintain prodⁿ grade K8s clusters from the command line.

Kubernetes Object Model:

- * Spec section: declare the desired state of the object.
- * Status section: K8s system manages the status section for objects. Records the actual state of the obj.
- * Objects: Pods, Replicasets, Deployments, Namespaces
- * 

```

graph TD
    A[Obj. defn (yAML)] --> B[Obj. defn (changed to JSON by Kubelet)]
    B --> C[API Server]
  
```

- * Pod:
 - The containers in a pod share the same network namespace, meaning that they share a single IP add. Originally assigned to the pod.
 - ephemeral in nature, and they do not have the capability to self-heal themselves. So, they are used with controllers which handle Pod's replication, fault tolerance, self-healing.
 - e.g. (controllers): Deployments, Replicasets, Replication Controllers

- * Labels:
 - key-value pairs attached to K8s objects
 - used to organize and select a subset of obj.
 - do not provide uniqueness to obj.
 - Controllers use labels to logically group together decoupled obj., rather than using objects' names or IDs.

- * Default controller \rightarrow Deployment which configures a ReplicaSet controller to manage Pod's lifecycle.

- *  With the help of a ReplicaSet, we can scale the number of pods running or specific app's container image.

- * Deployment:

- \rightarrow provides declarative updates to pods and replicasets.
- \rightarrow ensures that the current state always matches the desired state.
- \rightarrow directly manages its ReplicaSets for app's scaling.

- * Namespaces:

- \rightarrow Partition the K8's cluster into virtual sub-clusters called namespaces.

- \rightarrow `$ kubectl get namespaces` : list all namespaces

NAME	STATUS	AGE
default	Active	11h
Kube-public	Active	11h
Kube-system	Active	11h
Kube-node-lease	Active	11h

04 namespaces created by K8s

- \rightarrow Resource Quotas : Help users limit the overall resources consumed within Namespaces

- \rightarrow Limit Ranges : Help limit the resources consumed by pods or containers in a Namespace

Authentication, Authorization, Admission Control:

- * Authentication:

- \rightarrow K8s uses a series of authentication modules.

1. X509 client certificate
2. Static Token file
3. bootstrap tokens
4. Service account tokens
5. OpenID connect tokens
6. Webhook Token Auth
7. Authenticating Proxy

A Authorization:

- After a successful authentication, users can send the API req. to perform diff. operⁿs.

→ Authorization models:

Node → authorizes API req. made by kubelets

Attribute-based access control (ABAC)

Webhook - K8s can request authⁿ decisions to be made by 3rd-party services

RBAC -
a) Role → access to resources within a specific Namespace
b) ClusterRole → cluster-wide access

Services:

- * offers a single DNS entry for a containerized appⁿ managed by the K8s clusters, regardless of the no. of replicas, by providing a common load balancing access point to a set of pods logically grouped and managed by a controller such as a Deployment, Replicaset or DaemonSet.

- * Service: logically groups pods and defines a policy of any user or appⁿ wants to access them. This grouping is achieved by labels and Selectors. This service name is registered with the cluster's internal DNS service.

- * User/Client connects to the Service via its ClusterIP

↓
forward traffic to one of the pods attached to it

- (*) * Kube-proxy is responsible for implementing the service config in order to enable traffic routing to an exposed appⁿ running in pods.

- (*) * As services are the primary mode of commⁿ between containerized appⁿs managed by K8s, K8s supports 02 methods for discovering services.

- i) Environment Variables: As soon as the pod starts on any worker node, the kubelet daemon running on that node adds a set of env. variables in the pod for all active services.

- ii) DNS: DNS record for each service.

- * Access scope of the service is decided by the ServiceType property.

- * ClusterIP - default ServiceType
 - ↳ virtual IP add. received by a service.
 - ↳ used for communicating with the service and is accessible only from within the cluster.
- * NodePort ServiceType → High-port picked up from the range (30000 - 32767) is mapped to the respective service from all the worker nodes.
 - ↳ useful when we want to make our services accessible from the external world.
- * LoadBalancer ServiceType: Nodeport ClusterIP (→ automatically created \bigcirclearrowleft)
 - ↳ will only work if the underlying infra. supports the automatic creation of load balancers.
- * A Service mapped to ExternalIP address if it can route to one or more of the worker nodes.
- ExternalIP address → not managed by K8s.
- * ExternalName : special Service Type that has no Selectors and does not define any endpoints.
 - ↳ make externally configured services available to apps inside the cluster.

Kubernetes Volume Management:

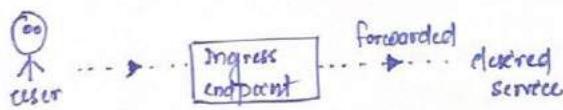
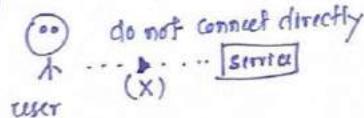
- * containers → ephemeral in nature. (all data inside a container is deleted if the container crashes).
- ↓
- * Volumes (storage abstractions)
- * Volume - mount point on the container's file system.
- * VolumeType - decides the properties of the directory, like size, content, access modes etc.
- e.g. emptyDir
hostPath
nfs
secret
configmap
PersistentVolumeClaim

- * Container Storage Interface (CSI) - designed to work on diff. container orchestrators
- ConfigMaps and Secrets:
 - * ConfigMaps allowed us to decouple the config details from the container image.
 - * Using ConfigMaps $\xrightarrow{\text{pass config. data as key-value pair}}$ Pools/sys. components \rightarrow in the form of env. variables
 - * Secrets $\xrightarrow{\text{share sensitive info (passwords, tokens, keys)}}$

Ingress:

- * Colln of rules that allow inbound conn's to reach the cluster services.
- ↓
configures a Layer 7 HTTP/HTTPS load balancer for Services

- * With ingress,



- * Ingress controller \rightarrow apply watching API server for changes in the Ingress resources and updates the Layer 7 load balancer.
 \Downarrow
 (K8s supports AWS, GCE, and NGINX Ingress controllers)

Quota and Limit Management:

- * Following types of quota per namespace:
 - Compute Resource Quota
 - Storage Resource Quota
 - Object Count Quota

Kubernetes Cluster Federation:

- * manage multiple K8s clusters from a single control plane.

Security Contexts:

- * Set discretionary access control for obj. access permissions
- * defines privilege and access control settings for a pod or container

Pod security policies: apply security settings to multiple pods and containers cluster wide.

Network Policies:

- * set of rules which define how ports are allocated to talk to other pods and resources inside and outside the cluster.
- * namespaced API resources

Deployments

- * Stateless app's
- * Pod dies random
reborn at the end.
- * pods deployed by Deployment are identical.

Stateful Sets

- * Stateful app's (databases)
* maintains pods deployed by stateful set are not identical.
- * Sticky identity for each pod
(fixed ordered names)
↓
persistent
identifiers
across any
rescheduling
- * When a pod dies
and replaced by a new one,
it maintains its identity.
- * all pods are not identical.
↓
Master
Slave
- * data persistence → recommended stateful app's.
If pods die, all data is gone.
→ data will survive, if pods die.
(Persistent Volume) → remote storage

mysq1-0
mysq1-1
mysq1-2

* fixed ordered DNS name

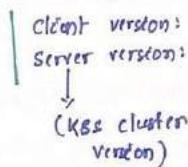
Objects:

- * pods
- { namespaces
- replicasets
- deployments
- statefulsets
- daemonsets
- Services
- configmaps
- volumes

- * To display all k8s objs → Kubectl api-resources

NAME	SHORTNAMES	KIND	APIGROUP	NAMESPACE
bindings		Binding		true
configmaps	cm	ConfigMap		true

- * To check k8s cluster's version → Kubectl version



Interacting with k8s:

- * To check the list of available API resources → Kubectl api-resources
- * To check the list of available API groups → Kubectl api-groups

admissionregistration.k8s.io/v1
admissionregistration.k8s.io/v1beta1

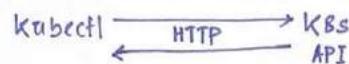
- * To list all pods in the namespace → Kubectl get pods
- * To list all pods in all namespaces → Kubectl get pods --all-namespaces
- * To list all pods in the current namespace with more details → Kubectl get pods -o yaml

- * Tools for testing k8s clusters and apps:

- 1) K8s dashboard
- 2) Kube-hunter
- 3) Project Quay
- 4) Kube-burner
- 5) Kube-bench
- 6) Kube-monkey

Kubectl:

- * Run interactive containers in your cluster \rightarrow Kubectl run
- * Kubectl logs $-f$
 - ↓
 - live stream new logs to your terminal
- * Cockpit to control K8s
- * Client for the K8s API

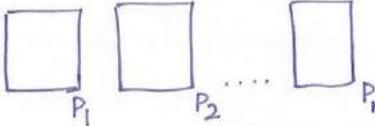


- * Supports JSONPath template:
 - ↓
 - {JSONPath expression}

Pod Concept:

- * A pod contains one or more appⁱⁿ containers which are relatively tightly coupled
- * A pod can contain init containers that run during pod start up.
- * Shared context of a pod is a set of Linux namespaces and cgroups.
- * Container image - nginx:1.14.2
obj. specfn & status - simple-pod.yaml
Create the pod: Kubectl apply -f https://k8s.io/examples/pods/simple-pod.yaml
- * Pods in a K8s cluster are used in 02 main ways:
 - 1) pods that run a single container (Pod - wrapper around a single container)
 - 2) pods that run multiple containers that need to work together

*  \longrightarrow meant to run a single instance of a given appⁱⁿ

 \longrightarrow Scaling your appⁱⁿ horizontally (to provide more overall resources by running more instances), you should use multiple pods, one for each instance.

Replication (in K8s)

- * Pods natively provide 02 kinds of shared resources for their constituent containers:

- 1) networking
- 2) storage

- * Pod remains on the node until the pod finishes execution.
- * A pod is not a process, but an env. for running containers.
- * You can use workload resources to create and manage multiple pods for you.

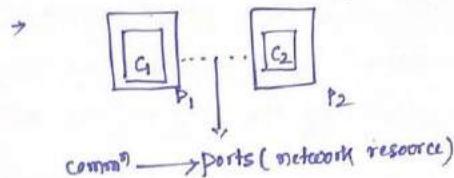
- 1) deployment
- 2) stateful set
- 3) daemonset

- * Pod templates → specify for creating pods and are included in workload resources such as deployments
jobs
daemonsets
- * Pod update operations → patch → partially update the specified pod
replace → replace the specified pod

- * Pod networking:

→ Pod assigned unique IP add.
→ Each container in a pod shares the same namespace including IP add. and network ports.

→ Inside a pod, the containers that belong to the pod can communicate with one another using local port.



- * Static pods:

→ managed directly by the kubelet daemon on a specific node without the API server observing them.

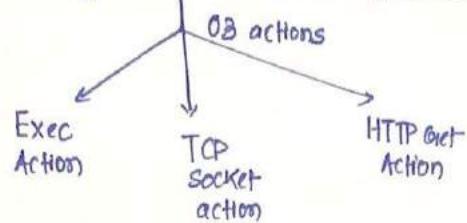
→ Kubelet directly supervises each static pod and restarts if it fails.

→ always bound to one kubelet on a specific node.

→ main use: Run a self-hosted control plane → using kubelet to supervise the individual control plane components

→ specify of a static pod cannot refer to other API objs

- * Probe: Diagnostic performed by the kubelet on the container



- * Pod evicted from the node due to lack of resources or, the node fails.
- * Kubectl get pods → inspect a pod running on your cluster

A pod can be in one of the following phases.

- 1) Pending → Pod has been created and accepted by the cluster, but one of its containers are not yet running.
- 2) Running
- 3) Succeeded → all containers in the pod have terminated successfully.
- 4) Failed → all containers in the pod have failed.
- 5) Unknown → state of the pod cannot be determined.

Sidecar Container:

- * Dedicated sidecar container for logging on an app pod
- * Examples: Monitoring agents, data loaders
- * Sidecar container must be placed on the same pod to use the same resources being used by the main container.
- * Sidecar containers do not provide a user-facing service.
- * run side-by-side with your main container.

Init Container:

- * Examples: Create user accounts
Perform DB migrations
Create DB schemas
- * Init containers are started and executed in sequence.
- * If a pod's init container fails, the kubelet repeatedly restarts that init container until it succeeds.
- * To specify an init container for a pod, add the `initContainers` field into the pod spec.
- * Init containers can contain utilities or custom code for setup that are not present in an app image.
- * During pod startup, the kubelet delays running init containers until the networking and storage are ready.

- * If the pod restarts or is restarted, all init containers must execute again.
- * A Pod cannot be in Ready status until all init containers have succeeded.
- * Init-containers can be given access to Secrets that app containers cannot access.
- * Provide an easy way to block or delay the startup of app containers until some set of preconditions are met.

Leveness Probe:

- * Kubelet uses liveness probes to know when to restart a container.
- * Apps running for long periods of time eventually transition to broken states, and cannot recover except by being restarted. K8s use liveness probe to detect and remedy such situations.

Kubectl describe pod liveness-exec
 Kubectl get pod liveness-exec → RESTARTS has been incremented.

Workload Objects:

ReplicaSet:

- * Purpose is to maintain a stable set of replica pods at any given time.
- * Identifies new pods to acquire by using its selector.
- * A deployment is a higher-level concept that manages ReplicaSets and provides declarative updates to pods.
- * Name of a ReplicaSet obj. must be a valid DNS subdomain name.
- * A ReplicaSet can be auto-scaled by an HPA (Horizontal Pod Scaler).
- * Deployments own and manage their ReplicaSets.
- * When load becomes too much for the no. of existing instances, K8s enables you to easily scale up your app, adding additional instances as needed.

Deployment:

- * define Deployments to create new Replica sets
- * create the deployment
 - Kubectl apply -f <https://k8s.io/examples/controllers/nginx-deployment.yaml>

Check whether the deployment was created

Kubectl get deployments

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
nginx-deployment	0/3	0	0	1s

How many replicas of the app are available to the users (ready/created)

kubectl rollout status deployment/nginx-deployment

waiting for rollout to finish: 2 out of 3 replicas have been updated.

kubectl get deployments

	NAME	READY	UP-TO-DATE	AVAILABLE	AGE
	nginx-deployment	3/3	3	3	18s

kubectl get rs

	NAME	DESIRED	CURRENT	READY	AGE
	nginx-deployment	3	3	3	18s

- * Pod-template-hash label $\xrightarrow{\text{added by}}$ Deployment controller to every replicaset that a deployment creates.

- * scale a deployment:

kubectl scale deployment/nginx-deployment --replicas=10

- * Rolling Update deployments support running multiple versions of an app at the same time.

- * Deployment status: $\xrightarrow{\text{progressing}} \text{complete}$
 $\xrightarrow{\text{fail to progress}}$

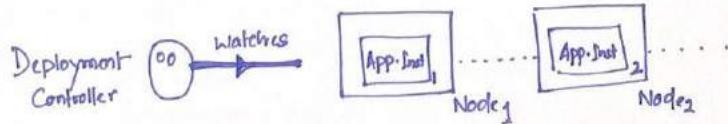
- * Canary deployment $\xrightarrow{\text{roll out releases to a subset of users or servers using the deployment.}}$

- * Strategy $\xrightarrow{\text{specifies the strategy to replace old pods by new ones.}}$

Recreate $\xrightarrow{\text{}} \text{Rolling update (default)}$

- * Deployment $\xrightarrow{\text{creating and updating instances of your app}}$

- * Deployment $\xrightarrow{\text{}} \text{API server} \xrightarrow{\text{}} \text{Kube-scheduler}$



Stateful Sets:

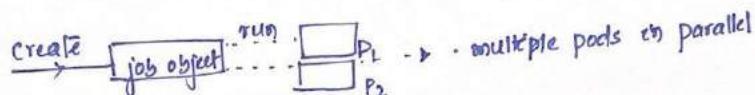
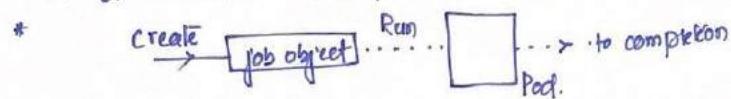
- * Provides guarantees about the ordering and uniqueness of pods.
- * Kind - Service

DaemonSet:

- * Daemonset pods are created and scheduled by the Daemonset controller.

Job:

- * Creates one or more pods and will continue to retry execution of the pods until a specified no. of them successfully terminate. When a specified no. of successful run of pods is completed, then the job is considered complete.



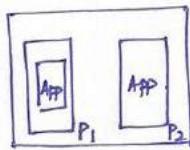
- * Jobs are complementary to Replication Controllers.
- * needs no pod selector
- * To make your job fail, you need to exit with a code other than 0.

Cron Job:

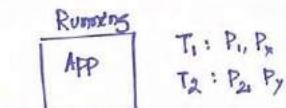
- * Timezone $\xrightarrow{\text{Set}}$ Kubernetes config.
- * Running backups or sending emails

Service:

- * An abstract way to expose an app running on a set of pods as a network service.

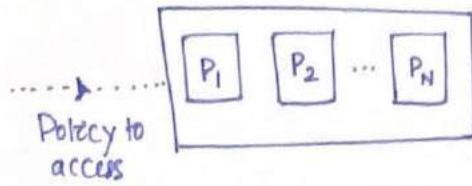


- * In a Deployment, a set of pods running in one namespace could be diff. from the set of pods running that app a moment later.



Problem

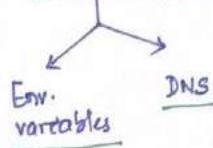
How do the frontends find out and keep track of which IP address to connect to so that the frontend can use the backend part of the workload?



Service → abstraction which defines a logical set of pods and a policy by which to access them.

- * set of pods targeted by a service is determined by a selector.

- * K8s supports 02 primary modes of finding a service.



- * Create a service:

Kubectl expose deployment/my-nginx → service/my-nginx exposed

- * Exposing a service onto an external IP add.

K8s supports it :
 1) nodeports
 2) loadbalancers

Volume & Storage Objects:

- * Provisioning of PVs

static
(cluster
admin)

dynamic

When none of the static PVs, the admin created match a user's PVC, then the cluster may try to dynamically provision a volume specially for the PVC. This provisioning is based on the StorageClasses.

- * Claim Ref: bi-directional binding between PV and PVC

- * PVs can have various reclaim policies:

1) Retain
 2) Recycle
 3) Delete (Default)

Autoscaling:

- * HPA → automatically updates a workload resource (deployment/statefulset) with the aim of automatically scaling the workload to match demand.

- HPA controls the scale of a Deployment and its Replicaset.
- K8s implements HPA as a control loop that runs intermittently.
- Common use: Fetch metrics from aggregated APIs
- While managing the scale of a group of replicas using the HPA, it is possible that the no. of replicas keep fluctuating frequently due to the dynamic nature of the metrics evaluated. (Known as thrashing or flapping) Similar to hysteresis in cybernetics

Taints and Tolerations:

- * Node Affinity is a property of pods that attracts them to a set of pods/nodes
- * Taints: allow a node to repel a set of pods.
- * Tolerations: are applied to pods and allow the pods to schedule onto nodes with matching taints.

Kube-proxy:

- * network proxy that runs on each node in your cluster, implementing part of the K8s service concept.
- * maintains network rules on nodes. These network rules allow network communication to your pods from network sources inside or outside of your cluster.
- * is only used to reach services

Namespace:

- * Kubectl api-resources --namespaced=true (to see K8s resources are there in a NS)
- * Kubectl api-resources --namespaced=false
- * K8s uses namespaces to organize objs in the cluster.

Kubeadmin:

- * tool built to provide Kubeadmin init and Kubeadmin join } for K8s clusters.
- * By design, it cares about bootstrapping, not about provisioning machines.

Kops: easiest way to get prod grade K8s cluster up and running, but also well provision necessary cloud infra.

Kubespray: deploy K8s clusters on multiple machines by combining Ansible and K8s.

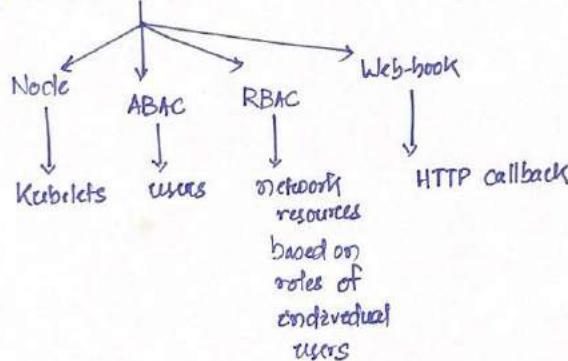
Controlling Access to the K8s API:

- * API Server on port 443, protected by TLS.
- * API Server serves HTTP on 02 ports.
 - I) localhost port - 8080
 - II) Secure port - 6443
- * Authentication: Service accounts are tied to a set of credentials stored as Secrets.

→ Bootstrap tokens are defined with a specific type of secrets that lives in the Kube-system NS.

* Authorization:

→ 04 authorization modes



* Admission Controller: piece of code that intercepts requests to the K8s API Server prior to persistence of the obj.

* Container-to-container communication: Shared volume in a pod
IPC namespace

* pod-to-pod comm: pods usually connect to other pods through a service.
Network plugin called Calico

* pod-to-service comm: Kube-proxy

* external-to-service comm: Ingress controllers

- KEDA: (Kubernetes-based Event Driven AutoScaler)



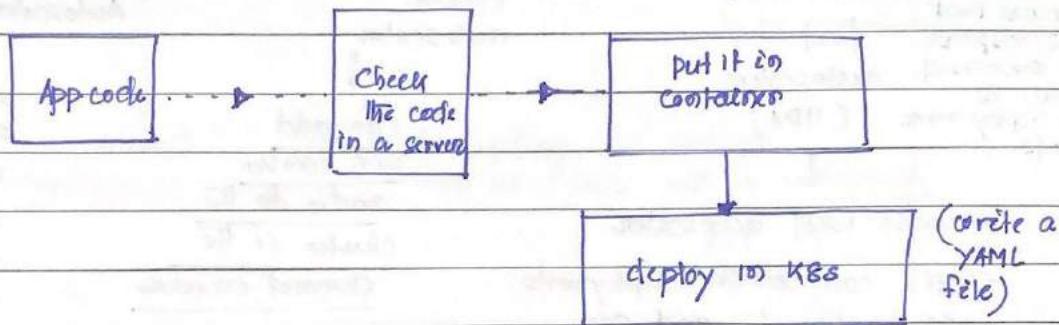
Proj. KEDA can be used to scale the k8s workload based on events triggered by external systems.

2019 (Microsoft - RedHat partnership)

- * K8s objects can be described in a data serialization lang. called yaml.
- * Main difference between ConfigMaps and Secrets

↓
Base64 encoding

(Cloud-Native Application Delivery)



- Application Delivery Fundamentals:

- * Git - standard version control system

↳ cheap local branching
staging areas
multiple workflows

- * IaC (Infra. as code) → process of managing and provisioning data centers through machine readable def'n files. rather than phy. H/W config.

↓
Ansible Autom8 Platform
Play books are used to describe the desired state of your infra. → config8/network policies/security
↓ (can be described as code)

- CI/CD:

CI: High automation and usage of version control

CD: autom. of deployment process

Popular CI/CD tools:

Spinnaker

GitLab

Jenkins

Jenkins X

Tekton CD

Argo

- GitOps:

↳ integrates the provisioning and change process of infra. with version control opns.

02 approaches

↓
Push-based

pipeline is started
and run tools
that make the changes
on the platform.

↓
Pull-based

an agent catches
the git repository
for changes and
compares the defⁿ
in the repo. with
the actual running state.
if changes are
detected, the agent
applies the changes
to the infra.

* K8s suited → GitOps (provides an API and
it's designed for
declarative provisioning
and change of
resources (right from
the beginning))

- Flux
 - ArgocD
 - ArgoCD (Kubernetes controller)
 - Flux (GitOps Toolkit)
- 02 popular
GitOps
frameworks

CLOUD NATIVE APPLICATION DELIVERY

GitOps:

- * declarative config is the key to dealing with infra. at scale.
 - * way to K8s cluster management and app deployment
 - * works by using Git as a single source of truth for declarative infra. and appms
 - * an operating model for K8s and other cloud native tech, providing a set of best practices that unify Git deployment, mgmt. and monitoring for containerized clusters and appms.
 - * Benefits:
 - a) increased productivity → continuous deployment autom (deploy fast)
 - b) enhanced developer exp. → push code and not containers
 - c) improved stability → convenient audit log of all cluster changes outside of K8s
 - d) higher reliability → Reduce your mean time to recovery (MTTR) from hours to mins (quick and easy recovery)
 - e) developer centric → deployment by pull request
 - * GitOps uses Git pull requests to automatically manage infra. provisioning and deployment.
 - * GitOps → an evolution to IaC that uses Git as the version control sys. for infra. configns
 - * In GitOps workflow, changes are made using pull requests which modify state in Git repository.
 - * GitOps gives you tools and a framework to take DevOps practices and apply them to infra. autom. and app deployment.
 - * GitOps: Running operations out of Git | by pull requests with the goal of continuous delivery of cloud native appn and typically Kubernetes.
- ↳ Tech.:
- a) IaC (Ansible, Terraform etc.) — Ops
 - b) Git (workflows) — Dev
 - c) CI/CD
 - d) convergent platforms (Kubernetes) — APIs, extensible

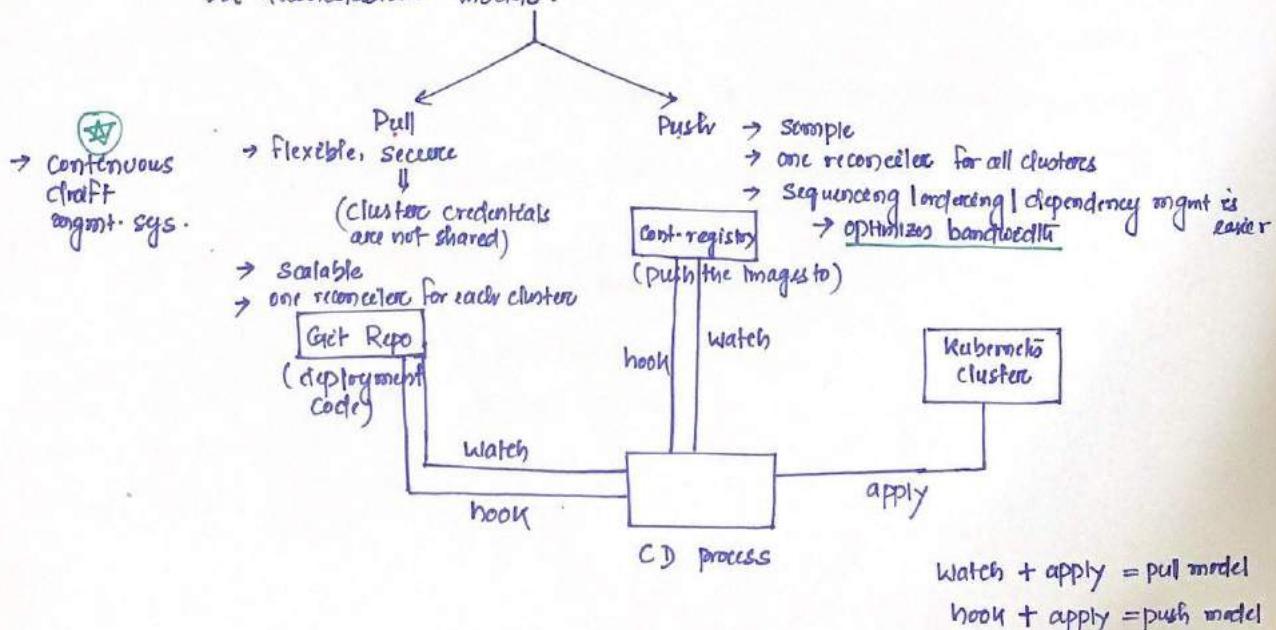
* Common usecases:

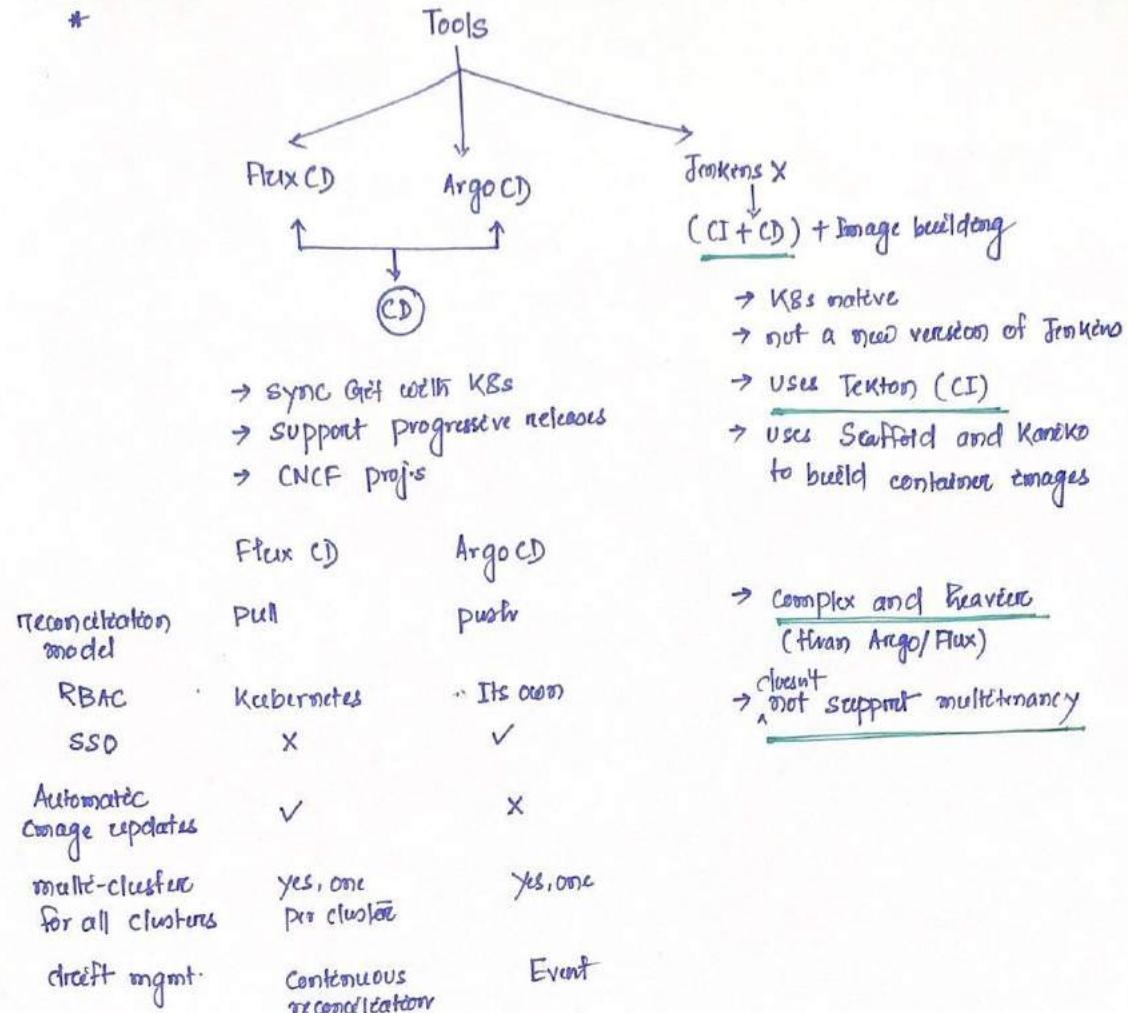
- a) continuous delivery of appl^{ing} configns (YAML manifest)
- b) apply release strategies (Blue Green
Rolling update
Canary)
- c) Infra. rollouts to k8s (ingress controllers
RBAC
network policies)
- d) Disaster Recovery
- e) Sync Secrets
- f) drift detection (identify and detect any changes in cluster config. settings)
- g) deploy to multiple K8 clusters
- h) securely handoff deployments to Engineering Teams.
- i) auto ^{update} k8s YAMLS on ^{new} Registry

* Principles and practices:

- a) ^{100%} configns → declarative code (YAML)
- b) Store desired state in Git
- c) apply approved changes automatically
- ④ d) check and correct configns with a s/w agent (reconciler)

* O2. Reconciliation models:





* WeaveWorkks → CNI Plugins (Weave)

↳ beyond CNI (CD automn) ... ➔ Flux ➔ GitOps (coined in 2016)

2016/2019 CSPs ➔ managed deployment services.

2019: Argo Flux (decided to merge and ~~later~~ changed it)

- ✓ Flux $\xrightarrow{\text{built}}$ GitOps Toolkit
- ✓ Argo $\xrightarrow{\text{built}}$ GitOps Engine

(Cloud Native Observability)

- Observability:

- * monitoring sub-category of observability
- * observability → closely related to the control theory which deals with behavior of dynamic systems.

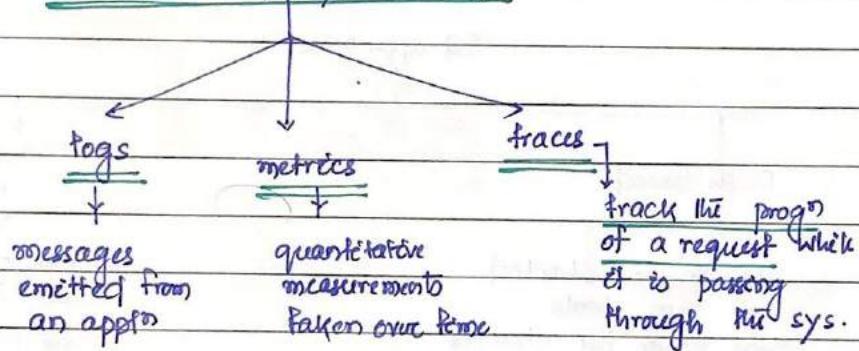
* You set the desired behavior of the sys. and trigger scaling events based on the load of the sys.

- * Goal: Allow analysis of the collected data.

↓
better understanding of the sys.
and react to error states

- Telemetry:

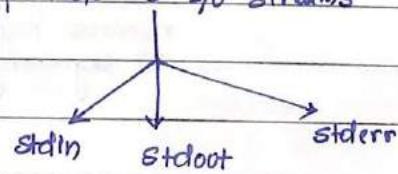
- * Container-based systems data



- Logging: levels →

debug
info
warning
error

- * Unix progs → provide 3 I/O streams



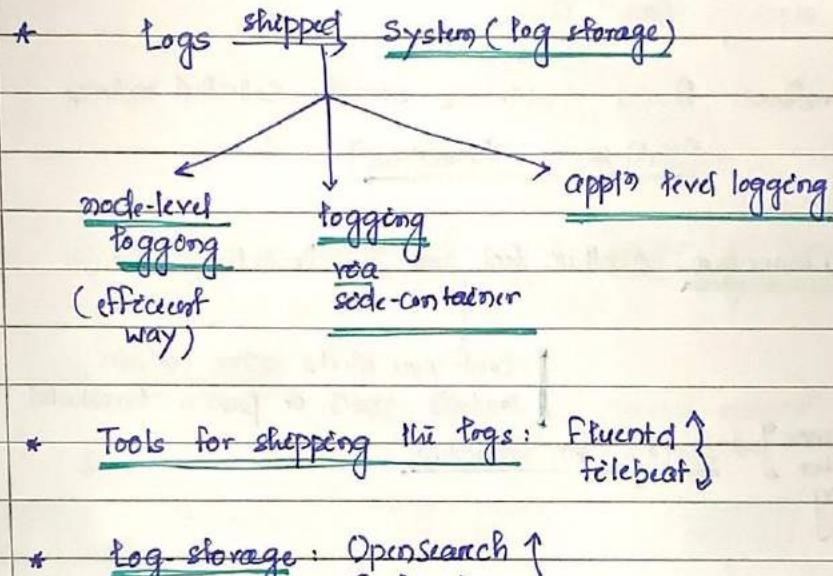
- * CLI tools : Docker, Kubernetes, to, show the logs

\$ docker logs nginx → container "nginx" logs

\$ kubectl logs nginx

\$ kubectl logs -p -c ruby web-1 → return snapshot of previously terminated ruby container logs from pod web-1

\$ kubectl logs -f -c ruby web-1 → stream logs in real-time (-f)

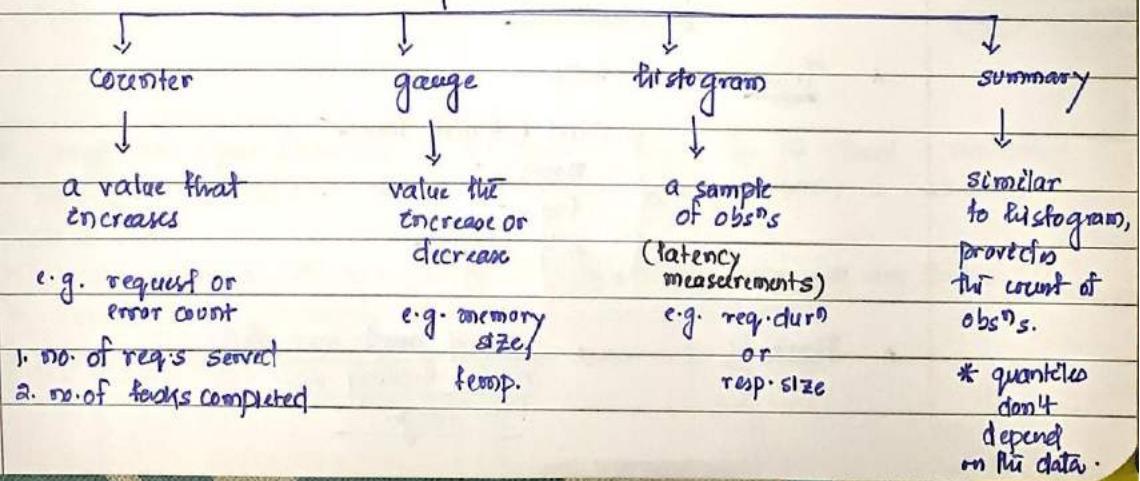


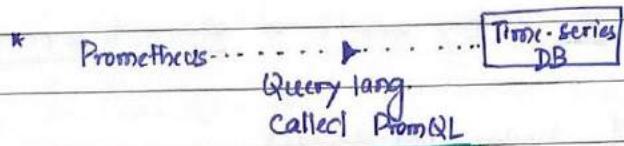
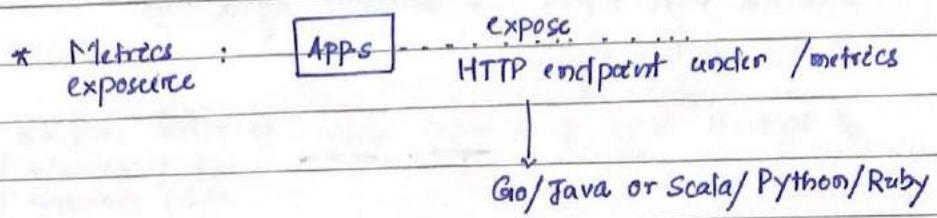
Prometheus:

* Open-source monitoring sys. → 2nd CNCF hosted proj. in 2016.

* Collects metrics emitted by apps and servers as time series data

* Prometheus data model provides 4 core metrics.

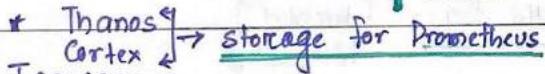




* Grafana: Build dashboards from the collected metrics.
[data source - Prometheus]

* Alert manager: another tool from the Prometheus ecosys.

↓
Configure alerts when certain metrics reach or pass a threshold

* Thanos 

Cortex
storage for Prometheus

- Tracing:

*

→ Microservice arch. ←

(How a req. is processed)

→ tracing
can be of
good use.

↓
describes the tracking of a req. where it passes through the services.

* Trace can include

↑ start & finish time ✓
 name ✓
 tag ✓
 tag msg ✓

* Traces → stored and analyzed in a tracing sys.
like Jaeger

* 2019: Open Tracing + OpenCensus → OpenTelemetry proj-
(CNCF proj.)

(projects)

↓
APIs,
SDKs,
Tools
used to
integrate
telemetry
such as metrics
protocols

* OpenTelemetry clients → Jaeger (platform) $\xrightarrow[\text{telemetry data}]{\text{export}}$

- Cost Management:

Automated and Manual optimizⁿ

↓
identify
wanted and
unused resources

↓
Right-sizing

↓
reserved
instances

↓
spot
instances

↓
auto-scaling helps
to shut down instances
that are not needed.

↓
on-demand
pricing model

↓
(reserve resources
and even pay
for them upfront)

↓
e.g. Batch jobs
or
heavy load
for short time

↓
get
unused
resources
First with
over-provisioned

FINOPS:

- * Cloud Fin. Mgmt (FinOps) is the operating model for the cloud - combining sys. best practices and culture to increase an org's ability to understand cloud costs.
- * FinOps empowers collaboration between IT, Engg., Finance, procurement and the biz.
- * Framework for managing operating expenditures in the cloud.
- * Financial accountability to the variable spend model of cloud.